



EECS 489

Computer Networks

Architecture and Applications

Agenda

- Application Requirements
- Network Layers
- A study of current design and choices

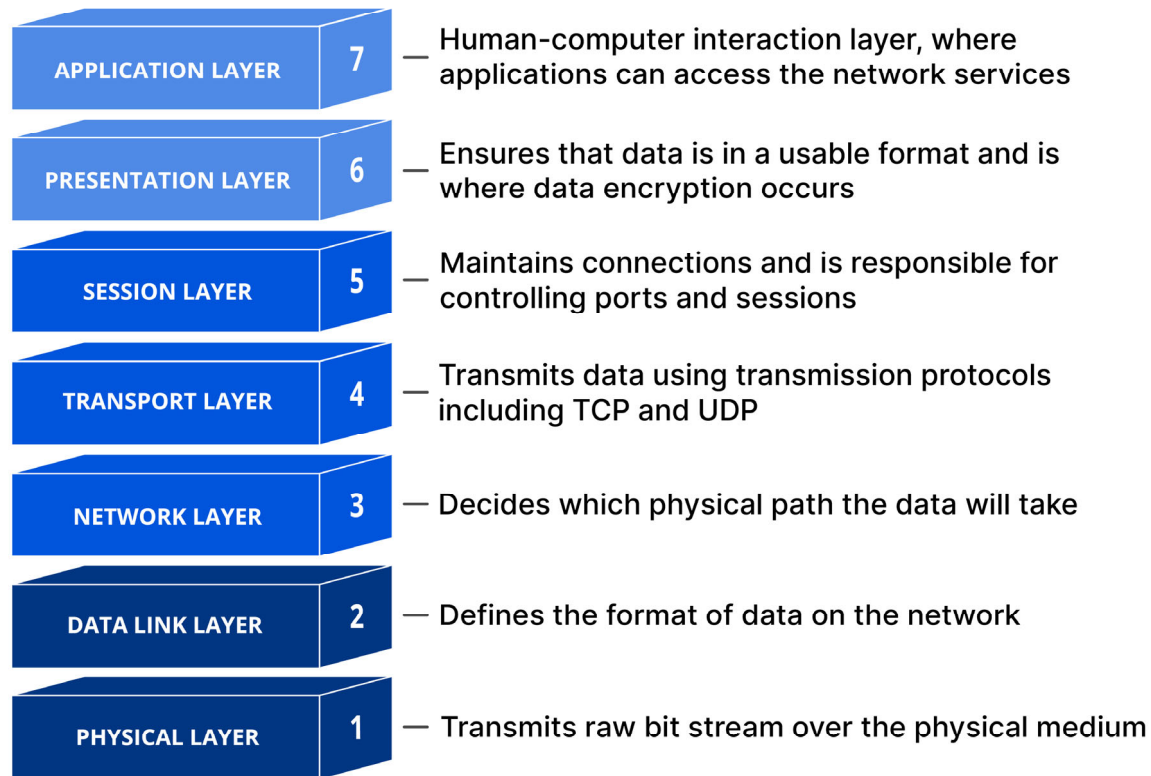
Announcements

- Anyone still on the waitlist please contact me after class
- Make sure you have access to:
 - AWS
 - The shared AMI for project 1
- Complete the mininet tutorial
- Take a look at the Ed post on Tips for using mininet

How to Design a Network


- Challenges
 - Has many users
 - Offers diverse services
 - Mixes diverse technology
 - Components build by different companies
 - Diverse ownership
 - Evolve over time

Review – OSI layers




OSI Communications

- Only Horizontal and Vertical Communications
- Each Layer offers a service to the higher layer using the services of the lower layer
- “Peer” layers on different systems communicate via a protocol

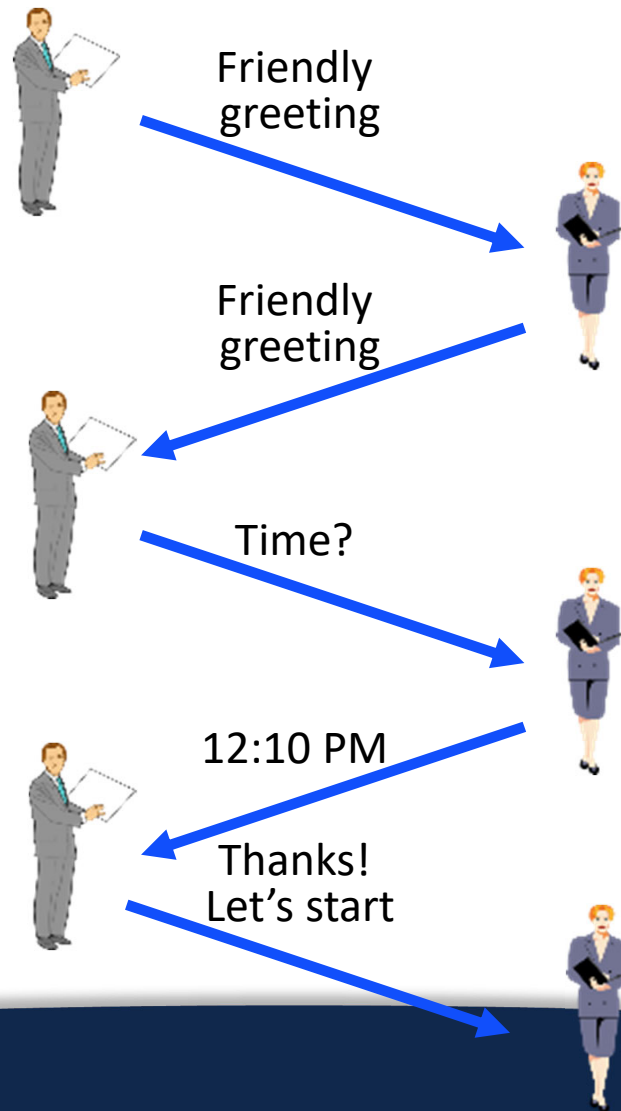


Layers and protocols

- Only Horizontal and Vertical Communications
- Each Layer offers a service to the higher layer using the services of the lower layer
- Communication between peer layers on different systems is defined by **protocols**



What is a Protocol?



What is a Protocol?

- An agreement between parties (in the same later) on how to communicate
- Defines the **syntax** of communication
 - **Header** → instructions on how to process **payload**
 - Each protocol defines the format of its headers
 - e.g., “the first 32 bits carry the destination address”

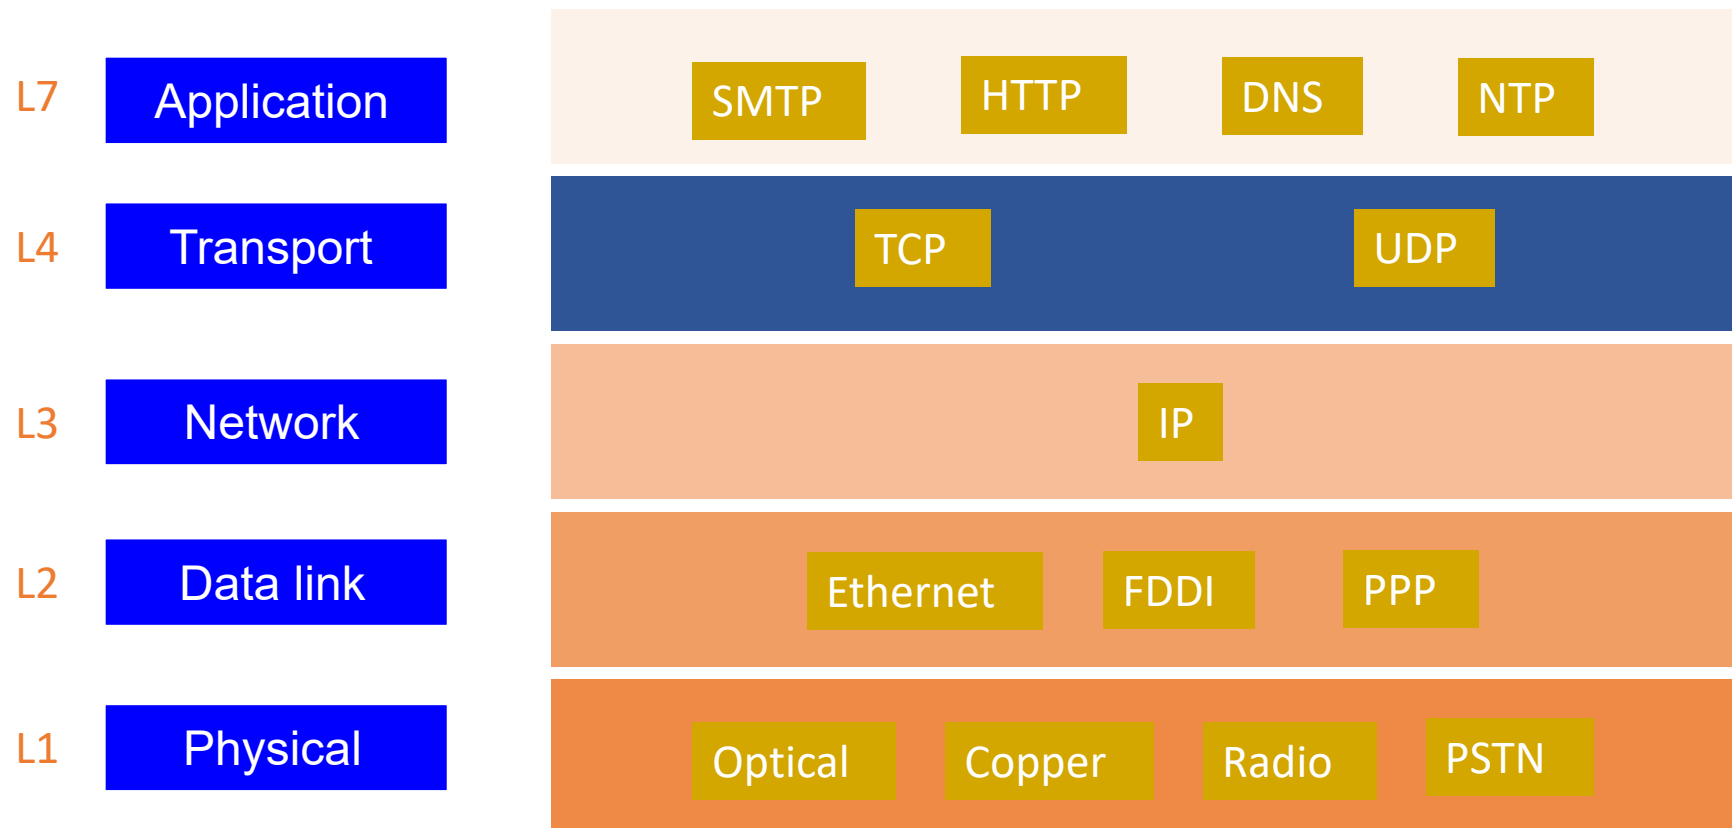
Header

Payload

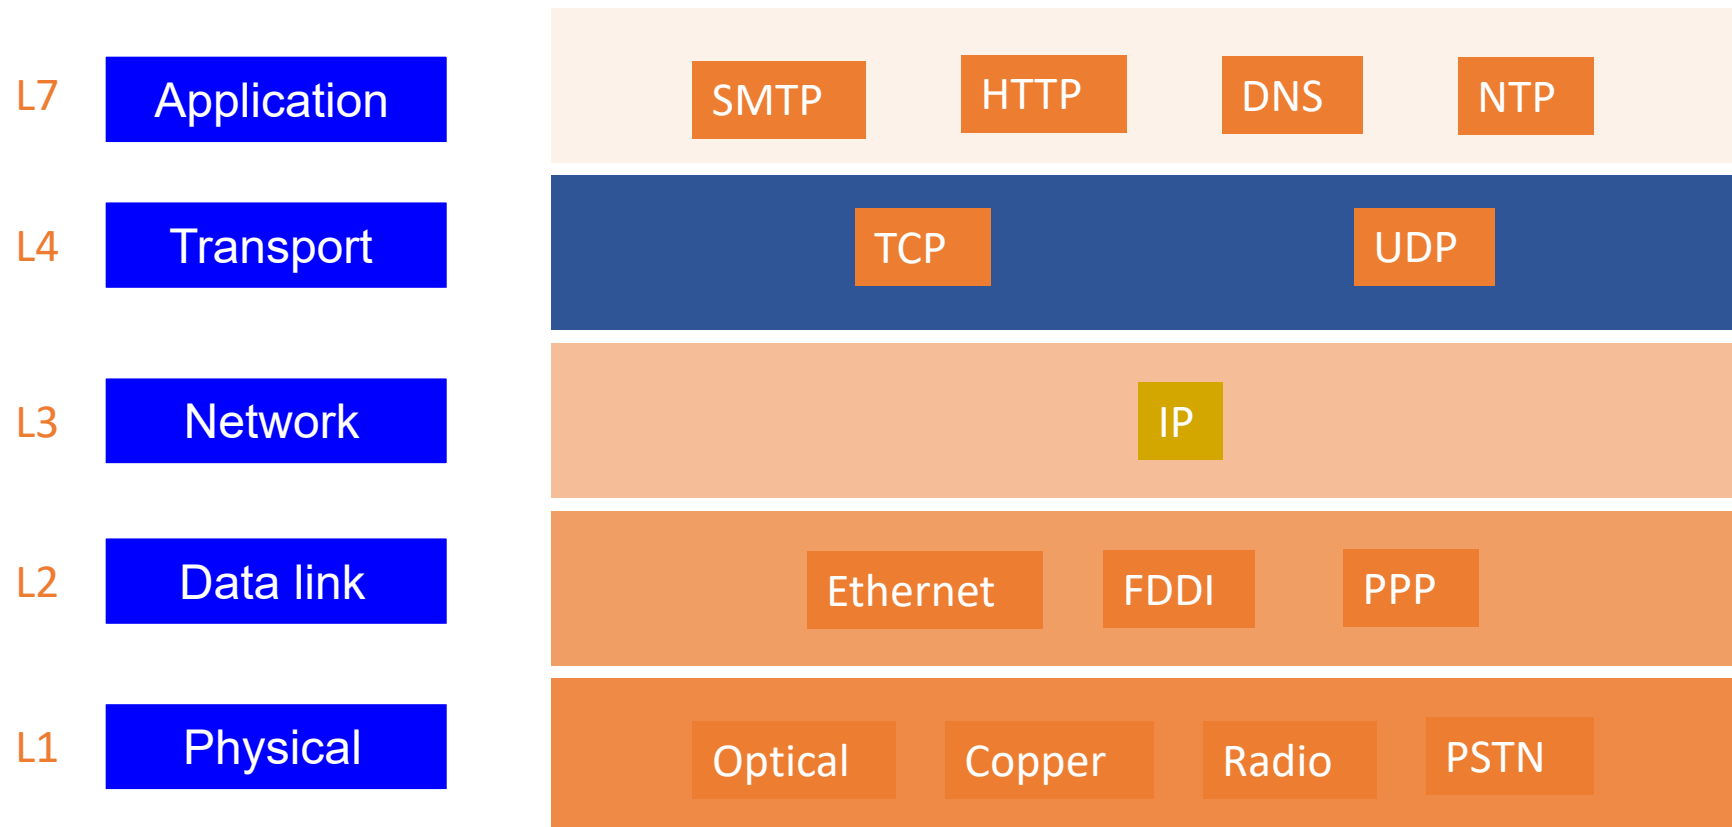
What is a Protocol?

- An agreement between parties on how to communicate
- Defines the **syntax** of communication
- And **semantics**
 - “First a hello, then a request...”
 - We will study many protocols later in the semester
- Protocols exist at many levels, hardware, and software
 - Defined by standards bodies like IETF, IEEE, ITU

Protocols at different layers




ONE network layer protocol



Starting out

- A Computer is switched on




- We need a scalable way to:
 - Get an IP address
 - Get IP address of the default gateway
- DHCP (Lecture 19 – Switched Networks)


The first packet



Who is communicating?



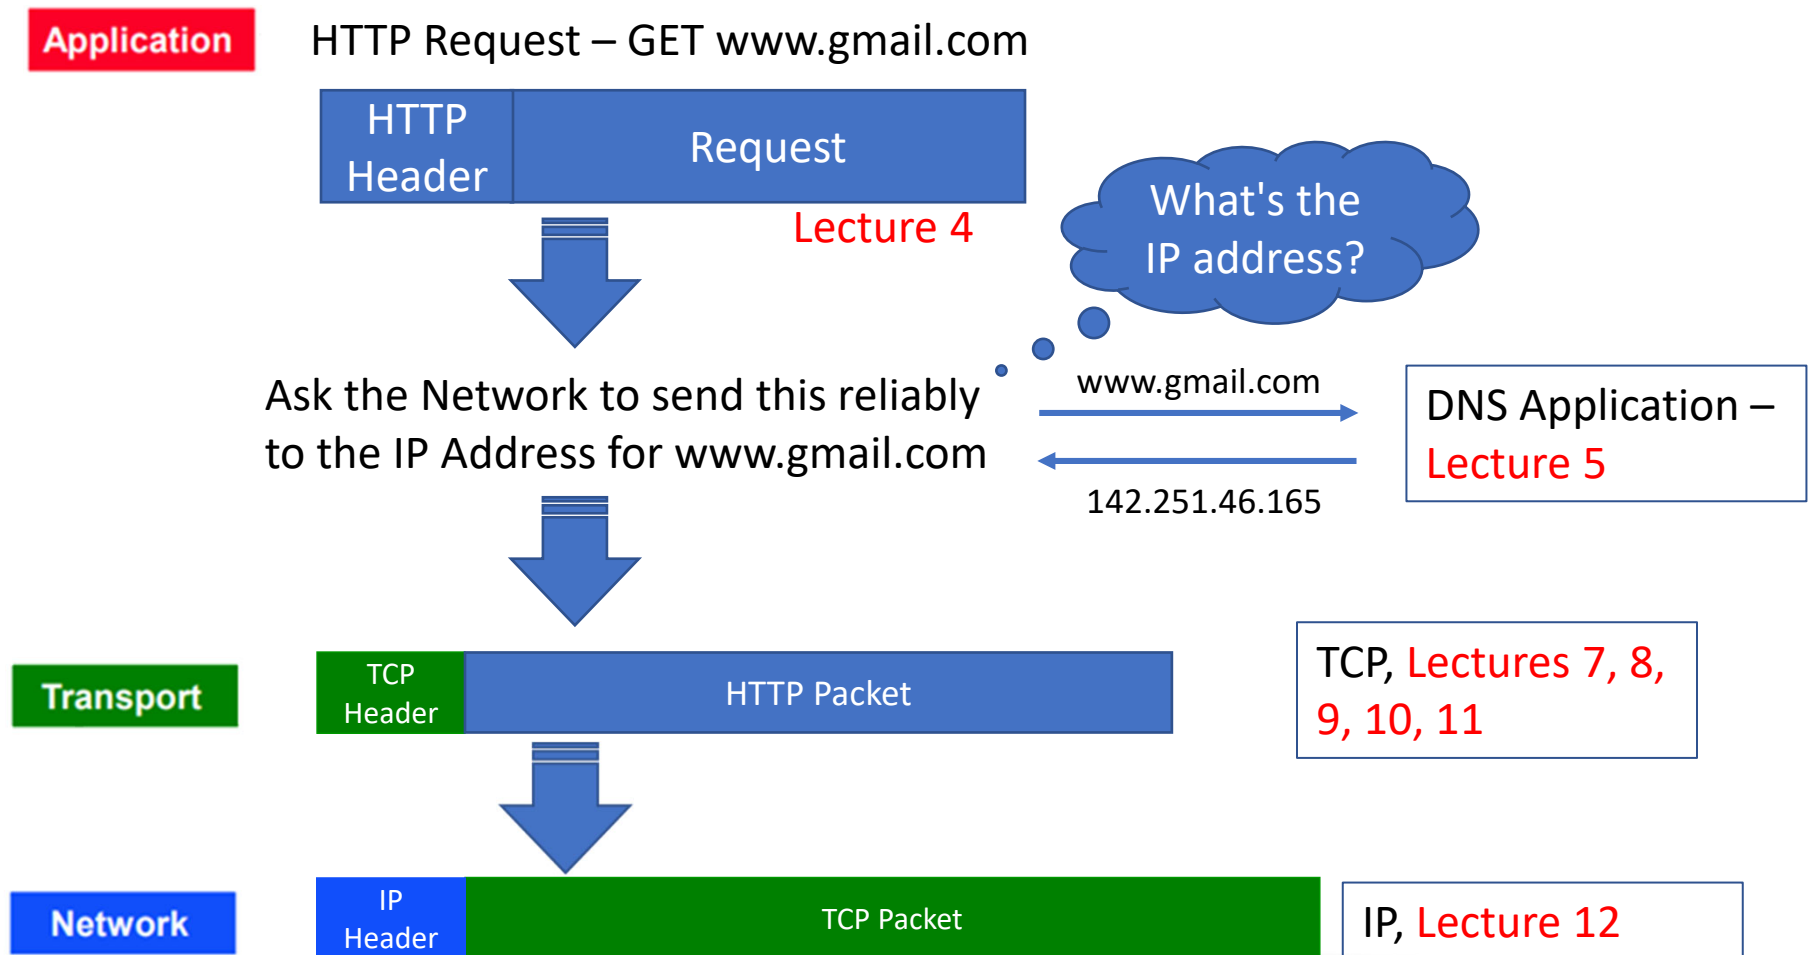
Actual Transfer of Data




Application Layer protocols

- Protocol between Web browser and Web server
 - HTTP – Lecture 4
- Other protocols?
 - FTP
 - SMTP
 - DNS
 - SSH
 - IMAP
 - SNMP
 - XMPP


Packet Generation



- How do we connect to the internet?

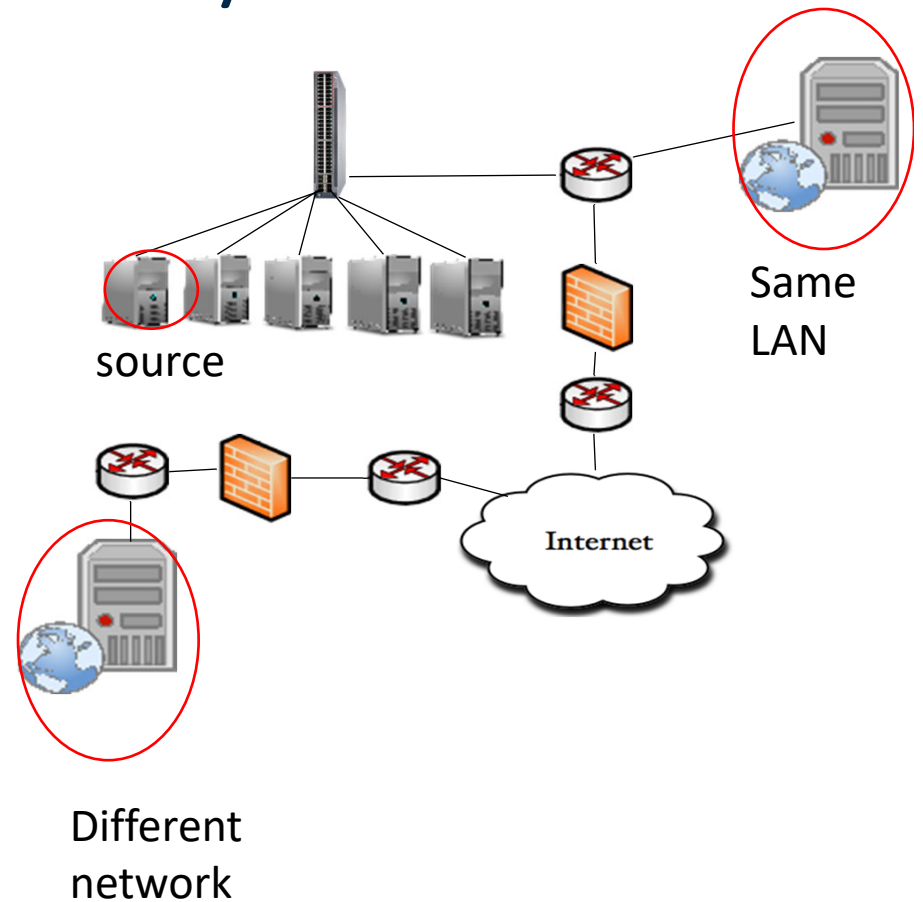


- More interesting case.

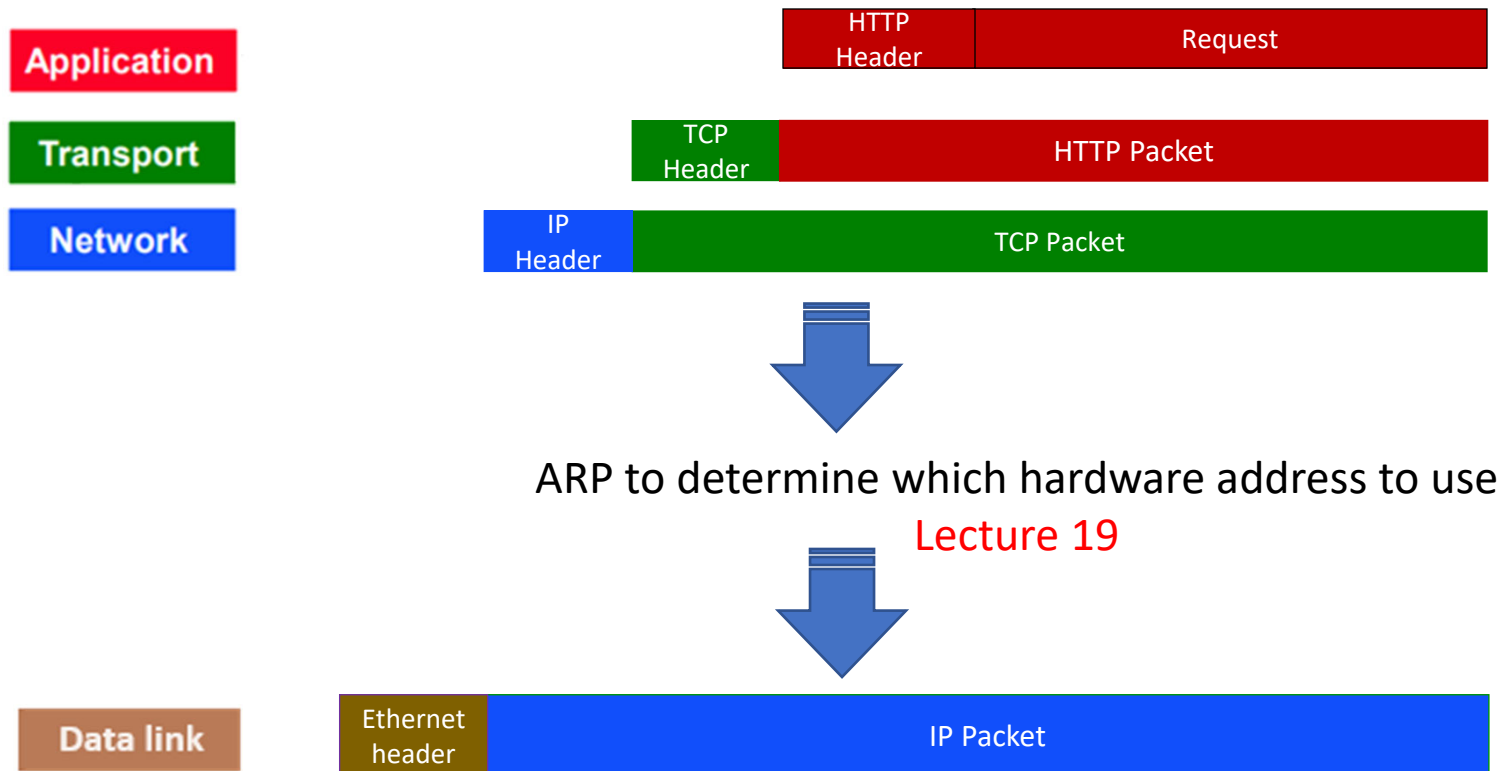


Packet Generation – IP Layer


- Where is the IP Address?
 - Same LAN
 - Different Network
- IP Addressing (**Lecture 16**)



Data Link Layer




What happens in the network




Routing and Forwarding
Lectures 14, 15, 16, and 17

At the destination

- What happens at the server?
- You are accessing gmail.com
- Option 1




- Option 2



Data Centers – Lectures 6 and 23

■ Tasks

- Is the user logged in
- Get user information
- Get user emails
- Get user calendar
- Is Free account?
 - Get ads for this user



Optimizations

- What can we optimize?
 - Routes
 - Content
 - URL and IP Address pairs
 - IP and MAC address pairs

What gets implemented at the end systems?

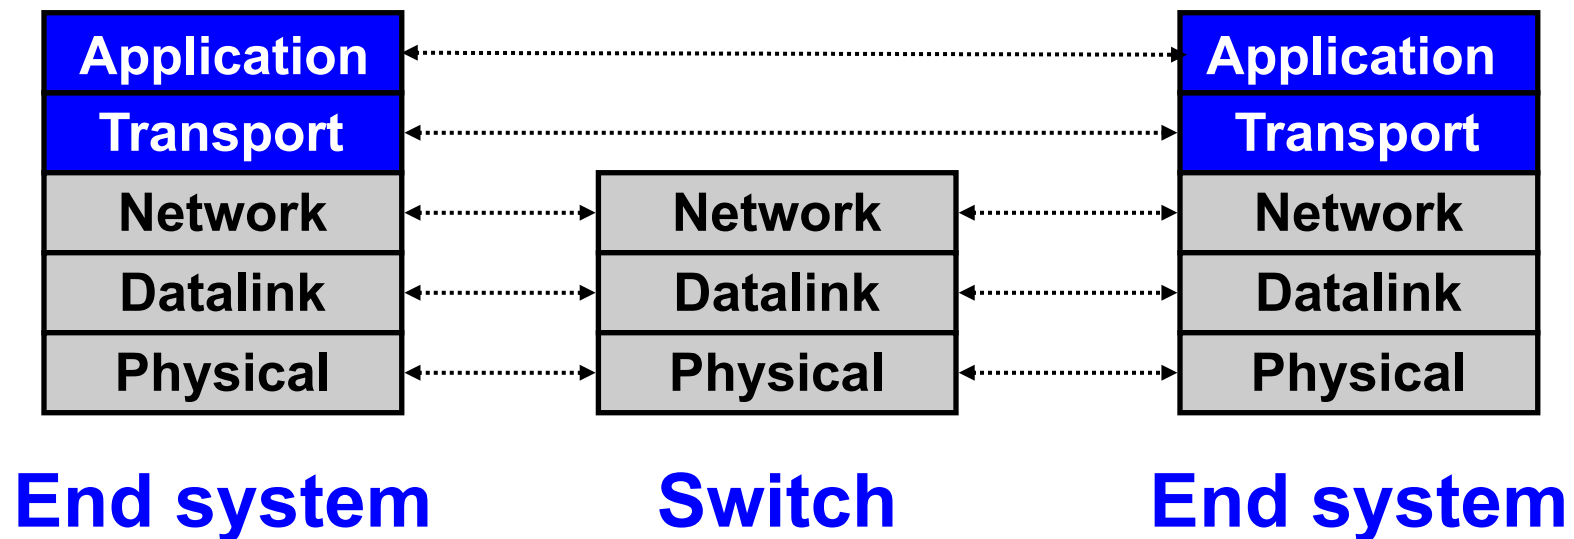
- Bits arrive on wire, must make it up to application
- Therefore, **all layers must exist at host!**

What gets implemented in the network?

- Bits arrive on wire → physical layer (L1)
- Packets must be delivered across links and local networks → datalink layer (L2)
- Packets must be delivered between networks for global delivery → network layer (L3)
- The network does not support reliable delivery
 - Transport layer (and above) not supported

Simple Diagram

- Lower three layers implemented everywhere
- Top two layers implemented only at hosts

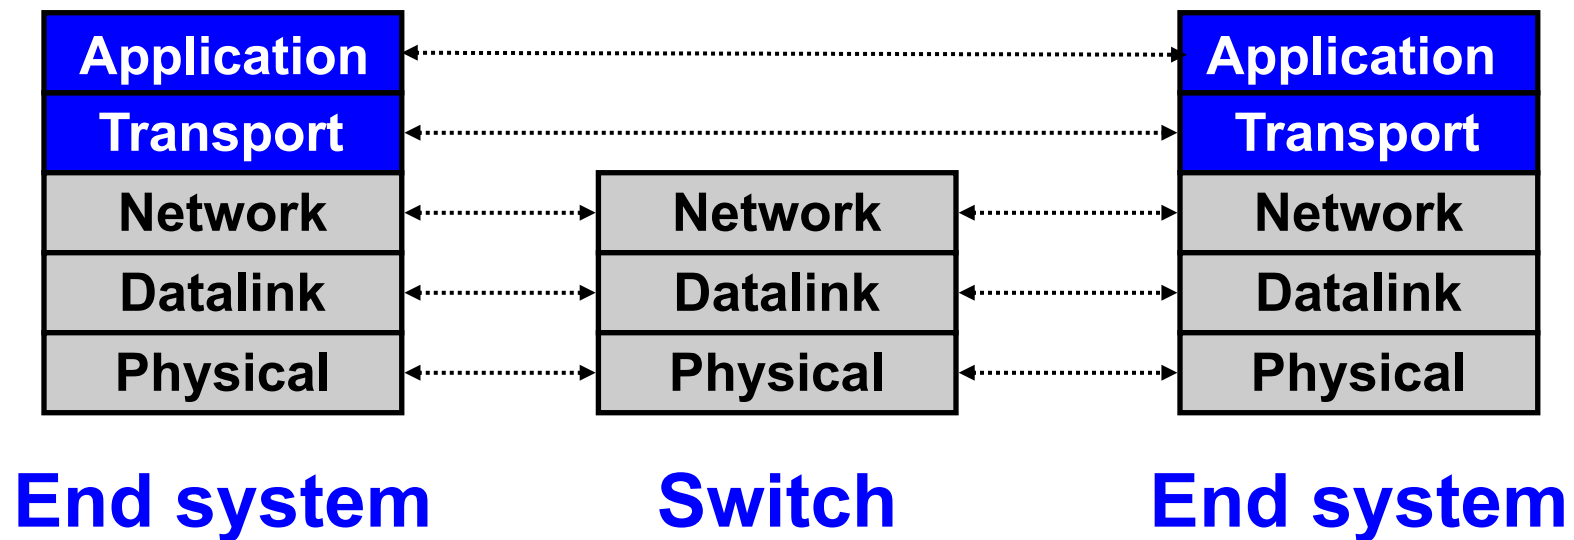


A closer look: End system

- Application
 - Web server, browser, mail, game
- Transport and network layer
 - typically part of the operating system
- Datalink and physical layer
 - hardware/firmware/drivers

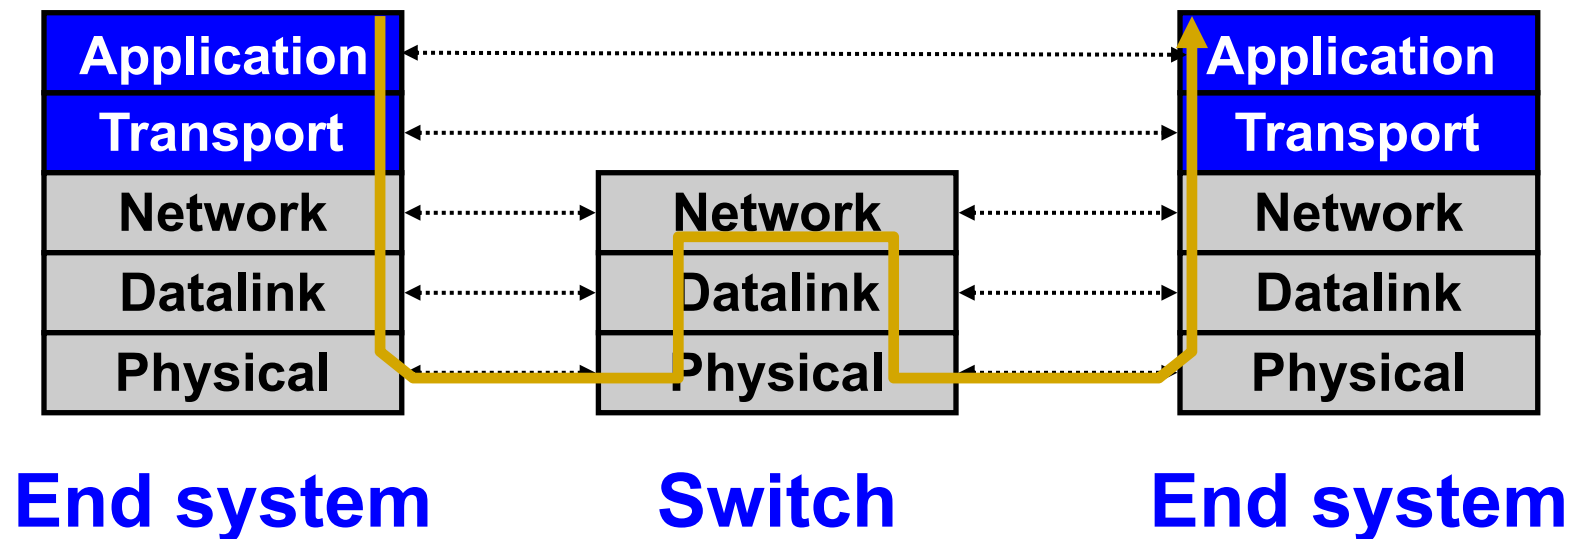
Logical communication

- A layer interact with its peers corresponding layer

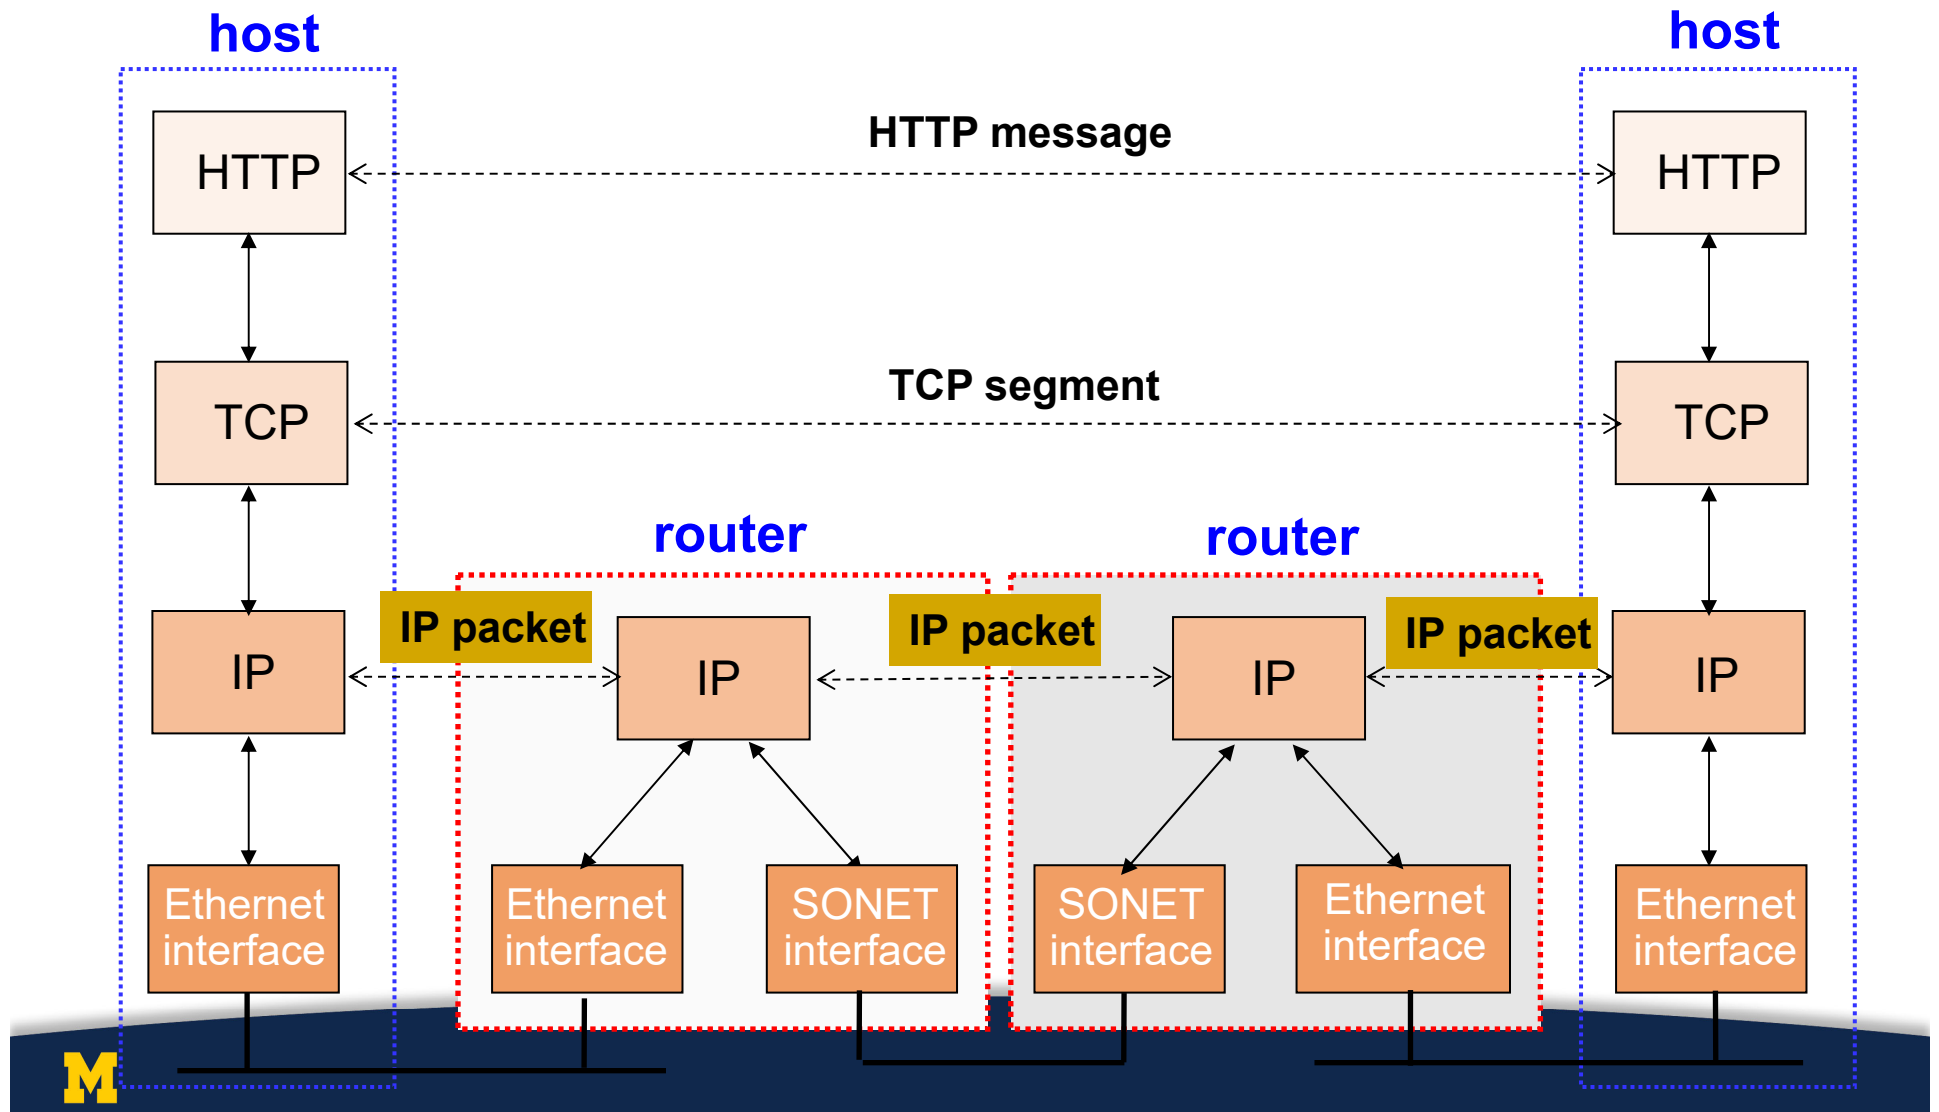


Physical communication

- Communication goes down to physical network
- Then up to relevant layer



A protocol-centric diagram



Pros and cons of layering

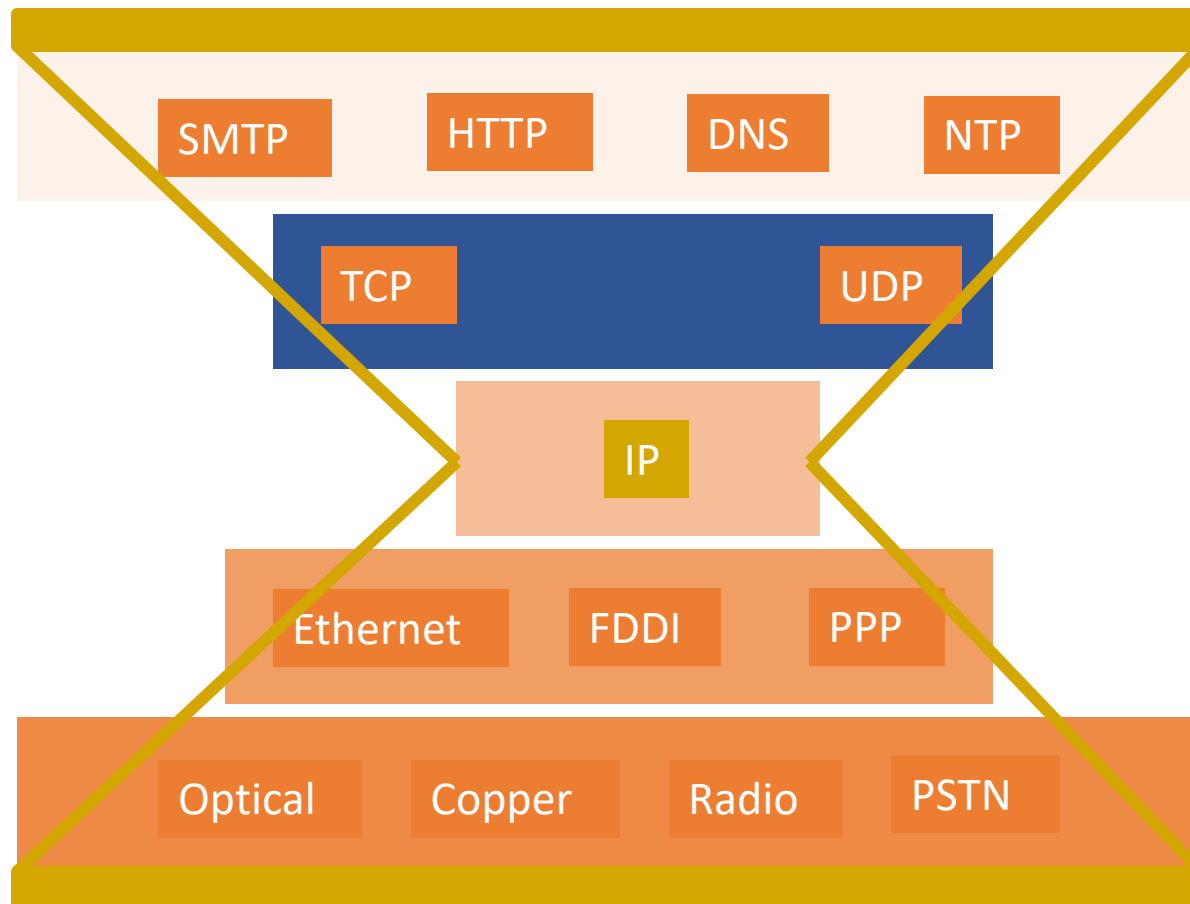
■ Why layers?

- Reduce complexity
- Improve flexibility

• Why not?

- Higher overheads
- Cross-layer information often useful

IP is the narrow waist of the layering hourglass



Implications of hourglass

- Single network-layer protocol (IP)
- Allows arbitrary networks to interoperate
 - Any network that supports IP can exchange packets
- **Decouples** applications from low-level networking technologies
 - Applications function on all networks
- Supports simultaneous innovations above and below IP
- But changing IP itself is hard (e.g., IPv4 → IPv6)

Placing network functionality

- **End-to-end arguments** by Saltzer, Reed, and Clark
 - Dumb network and smart end systems
 - Functions that can be *completely* and *correctly* implemented *only* with the knowledge of application end host, should not be pushed into the network
 - Sometimes necessary to break this for performance and policy optimizations
 - **Fate sharing**: fail together or don't fail at all

Internet Architecture

- Fundamental goal:
 - Effective network interconnection
- Goals, in order of priority:
 1. Continue despite loss of networks or gateways
 2. Support multiple types of communication service
 3. Accommodate a variety of networks
 4. Permit distributed management of Internet resources
 5. Cost effective
 6. Host attachment should be easy
 7. Resource accountability

Survivability

- If network disrupted and reconfigured
 - Communicating entities should not care!
 - No higher-level state reconfiguration
 - Ergo, transport interface only knows “working” and “not working.” Not working == complete partition.
- How to achieve such reliability?
 - Where can communication state be stored?

	Network	Host
Failure Handling	Replication	“Fate Sharing”
Network Engineering	Tough	Simple
Switches	Maintain State	Stateless
Host Trust	Less	More

Fate Sharing



- Lose state information for an entity if (and only if?) the entity itself is lost.
 - Example:
 - OK to lose TCP state if one endpoint crashes -
 - NOT okay to lose if an intermediate router reboots
 - Is this still true in today's network?
 - NATs and firewalls (More about those later)
- Survivability compromise: Heterogenous network -> less information available to end hosts and Internet level recovery mechanisms

Quiz 2 – open until 8pm

- <https://forms.gle/fqMfo2fDFHkBkADz7>



Types of Service

- Recall TCP vs. UDP
 - Elastic (in terms of delay) apps that need reliability: remote login or email
 - Inelastic, loss-tolerant apps: real-time voice or video
 - Others in between, or with stronger requirements
 - Biggest cause of delay variation: reliable delivery
 - Today's net: ~100ms RTT
 - Reliable delivery can add seconds.
- Original Internet model: “TCP/IP” one layer
 - First app was remote login...
 - But then came debugging, voice, etc.
 - These differences caused the layer split, added UDP
- No QoS support assumed from below
 - In fact, some underlying nets only supported reliable delivery
 - Made Internet datagram service less useful!
 - Hard to implement without network support
 - QoS is an ongoing debate...

Varieties of Networks

- Different types of networks
 - Interconnect the ARPANET, X.25 networks, LANs, satellite networks, packet networks, serial links...
- Minimum set of assumptions for underlying net
 - Minimum packet size
 - Reasonable delivery odds, but not 100%
 - Some form of addressing unless point to point
- Important non-assumptions:
 - Perfect reliability
 - Broadcast, multicast
 - Priority handling of traffic
 - Internal knowledge of delays, speeds, failures, etc.
- Much engineering then only has to be done once

Other goals

- Management
 - Today's Internet is decentralized - BGP
 - Very coarse tools. Still in the “assembly language” stage
- Cost effectiveness
 - Economies of scale won out
 - Internet cheaper than most dedicated networks
 - Packet overhead less important by the year
- Attaching a host
 - Not awful; DHCP and related autoconfiguration technologies helping. A ways to go, but the path is there

Accountability

- Huge problem.
- Accounting
 - Billing? (mostly flat-rate. But phones are moving that way too - people like it!)
 - Inter-provider payments
 - Hornet's nest. Complicated. Political. Hard.
- Accountability and security
 - Huge problem.
 - Worms, viruses, etc.
 - Partly a host problem. But hosts very trusted.
 - Authentication
 - Purely optional. Many philosophical issues of privacy vs. security.

Where do we go from here

- We start top-down and look at design choices, trade-offs, etc. at each layer