



EECS 489

Computer Networks

CDN and DNS

Announcements

- Please form teams for Assignment 2, 3, and 4
- Working without a group?
- Late days
 - 3 late days for the duration of the semester



Agenda

- CDN: Content Distribution Network
- DNS: Domain Name System

Recap: Improving HTTP performance

- Optimizing connections using three “P”s
 - Persistent (latency)
 - Parallel/concurrent (bandwidth and latency)
 - Pipelined over the same connection (latency)
- Caching
 - Forward proxy: close to clients
 - Reverse proxy: close to servers
- Replication

Replication

- Replicate popular Websites across many machines
 - Spreads load across servers
 - Places content closer to clients
 - Helps when content isn't cacheable

Content Distribution Networks (CDN)

- Caching and replication as a service
- Large-scale distributed storage infrastructure (usually) administered by one entity
 - e.g., Akamai is in 130 countries and 1700 networks
- Combination of caching and replication
 - **Pull**: Direct result of clients' requests (caching)
 - **Push**: Expectation of high access rate (replication)

Cost-effective content delivery

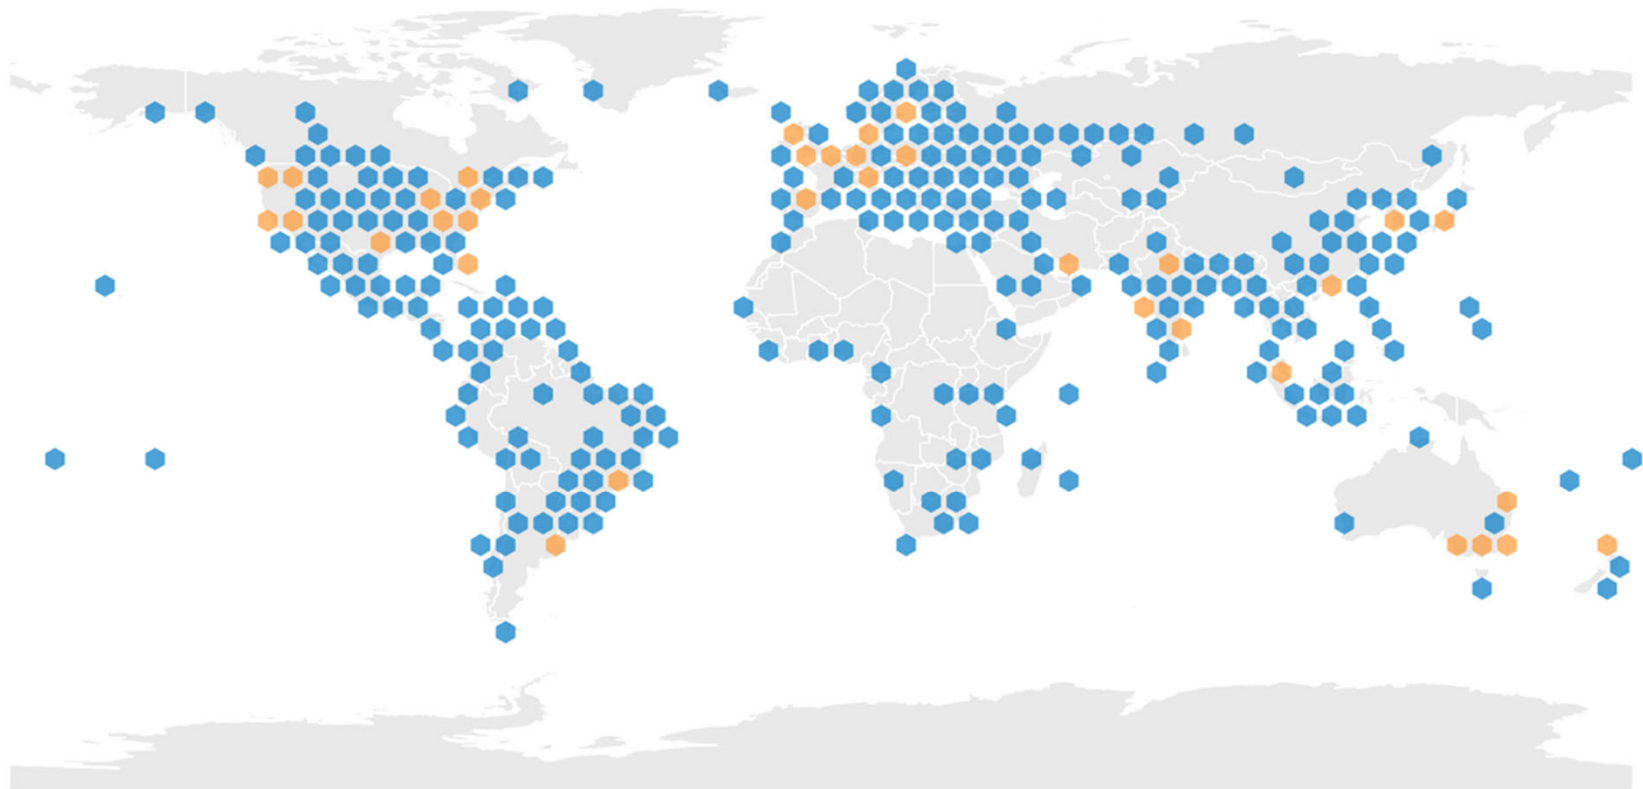
- General theme: multiple sites hosted on shared physical infrastructure
 - Efficiency of statistical multiplexing
 - Economies of scale (volume pricing, etc.)
 - Amortization of human operator costs
- Examples:
 - CDNs
 - Web hosting companies
 - Cloud infrastructure

CDN example – Akamai

- Akamai creates new domain names for each client
 - e.g., a128.g.akamai.net for cnn.com
- The client content provider modifies content so that embedded URLs reference new domains
 - “Akamaize” content
 - e.g., <http://www.cnn.com/image-of-the-day.gif> becomes <http://a128.g.akamai.net/image-of-the-day.gif>
- Requests now sent to CDN’s infrastructure

CDN example – Akamai

■ Akamai Media Delivery Network ■ Akamai Media Delivery + Storage



Retrieved on Sep 10, 2021 from <https://www.akamai.com/visualizations/media-delivery-network-map>

Why direct clients to particular replicas?

- Balancing load across server replicas
- Pairing clients with nearby servers to decrease latency and overall bandwidth usage

Packet Generation



DNS: Domain name system

Internet names & addresses

- Machine addresses: e.g., 141.212.113.143
 - Router-usable labels for machines
 - Conforms to network structure (the “where”)
- Machine names: e.g., cse.umich.edu
 - Human-usable labels for machines
 - Conforms to organizational structure (the “who”)
- The Domain Name System (DNS) is how we map from one to the other
 - A [directory](#) service

Why?

- Convenience
 - Easier to remember
- Provides a **level of indirection!**
 - Decoupled names from addresses
 - Many uses beyond just naming a specific host

DNS: History

- Initially all host-address mappings were in a `hosts.txt` file (in `/etc/hosts`):
 - Maintained by the Stanford Research Institute (SRI)
 - Changes were submitted by email and updates downloaded periodically from SRI
- As the Internet grew SRI could not handle load
 - Names were not unique anymore
 - Hosts had inaccurate copies of `hosts.txt`

DNS: History

- In 1983, the first stable operational DNS implementation included
 - The associated query protocol;
 - A server implementation; and
 - Initial root servers.
- Since inception, DNS scaled from 1000s of queries/day to 10s of billions queries/day

Goals

- Uniqueness: no naming conflicts
- Scalable
 - Many names and frequent updates (secondary)
- Distributed, autonomous administration
 - Ability to update my own (machines') names
 - Don't have to track everybody's updates
- Highly available
- Lookups are fast
- Perfect consistency is a **non-goal**

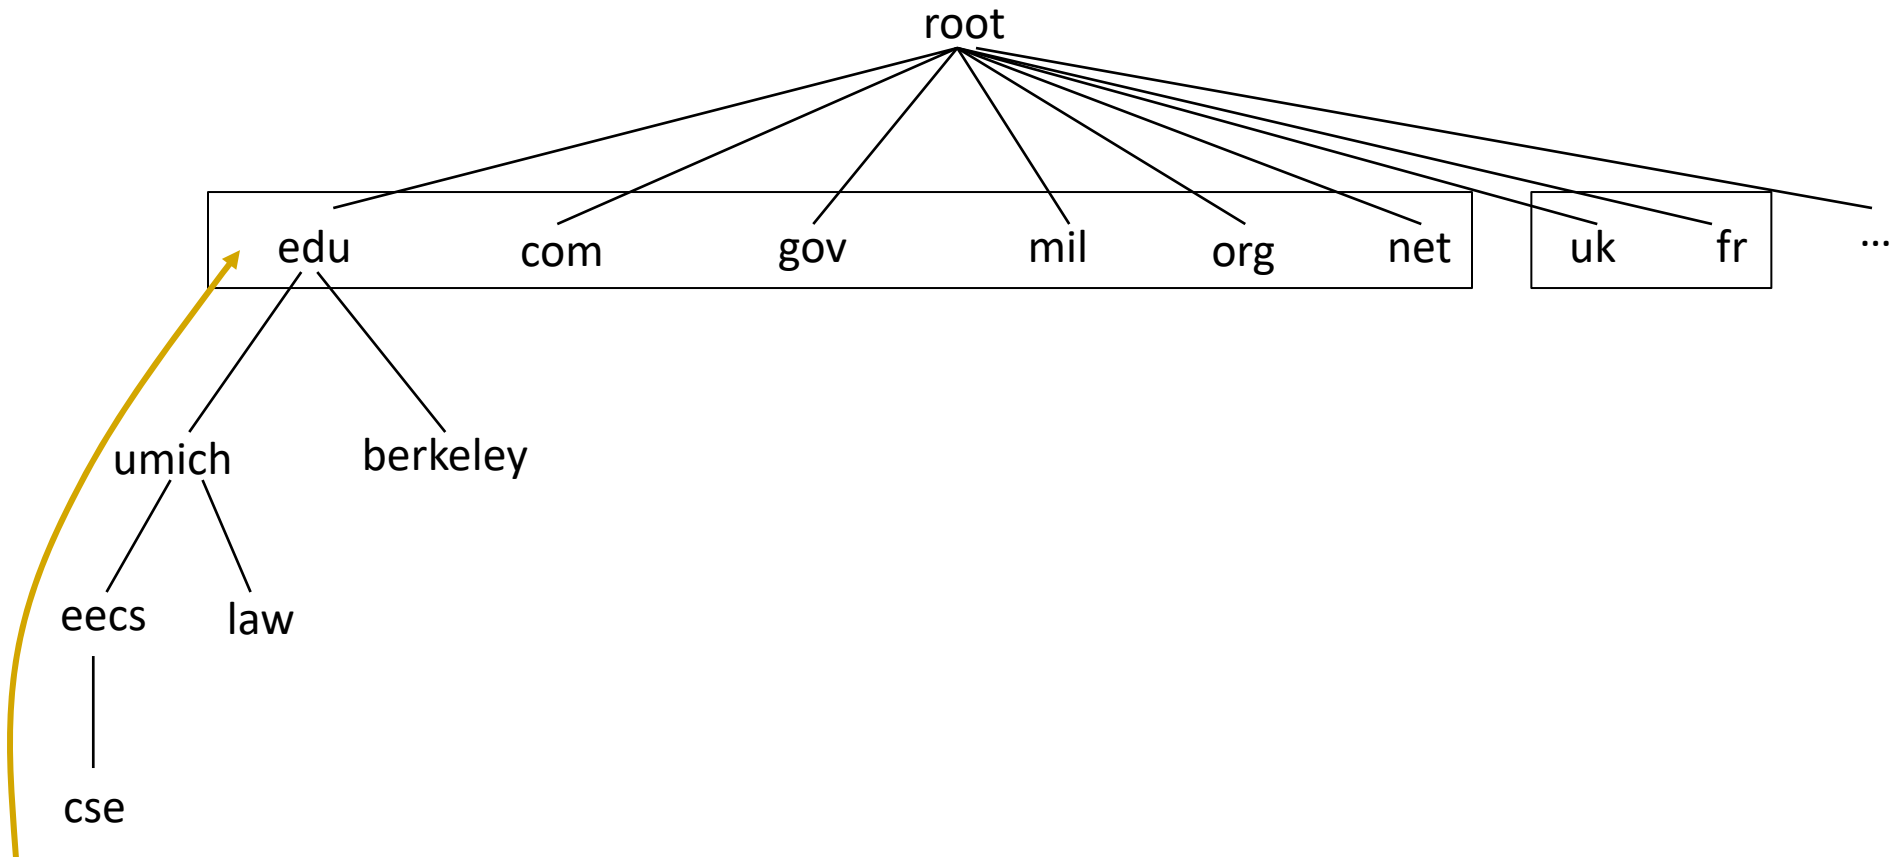
How?

- Partition the namespace
- Distribute administration of each partition
 - Autonomy to update my own (machines') names
 - Don't have to track everybody's updates
- Distribute name resolution for each partition
- How should we partition things?

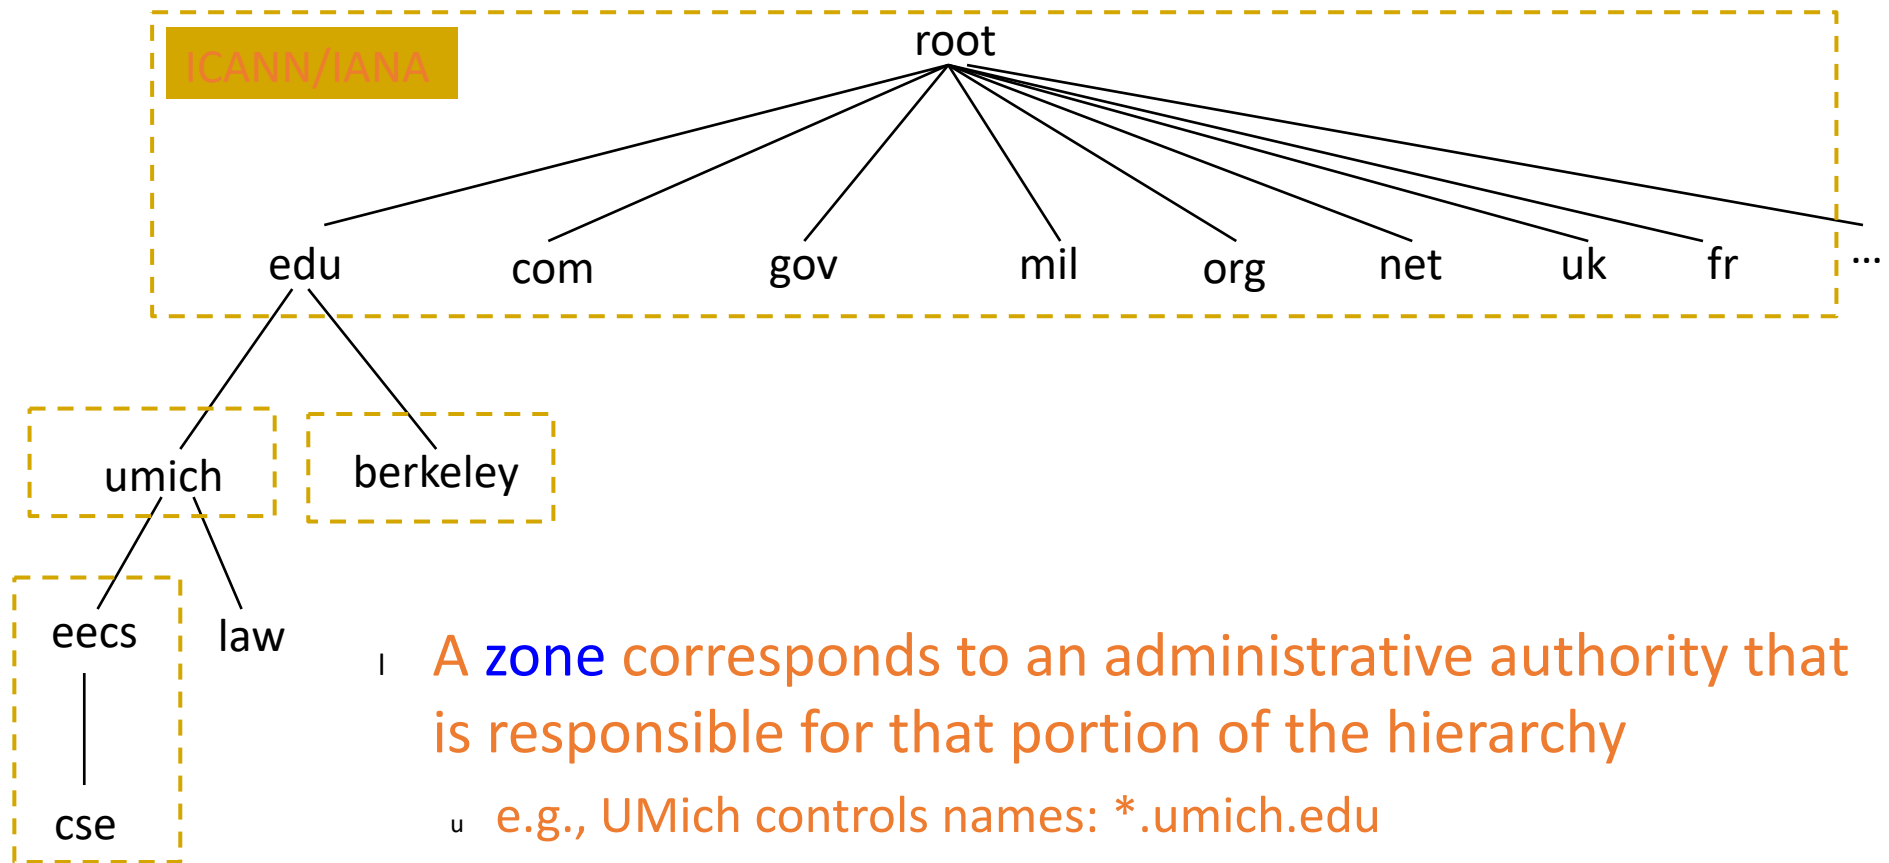
Key idea: Hierarchy

- Three intertwined hierarchies
 - Hierarchical namespace
 - As opposed to original flat namespace
 - Hierarchically administered
 - As opposed to centralized
 - (Distributed) hierarchy of servers
 - As opposed to centralized storage

Hierarchical namespace



Hierarchical administration



| A **zone** corresponds to an administrative authority that is responsible for that portion of the hierarchy

u e.g., UMich controls names: *.umich.edu

u e.g., EECS controls names: *.eecs.umich.edu

Server hierarchy

- Top of hierarchy: **Root servers**
 - Location hardwired into other servers
- Next Level: **Top-level domain (TLD) servers**
 - .com, .edu, etc.
 - Managed professionally
- Bottom Level: **Authoritative DNS servers**
 - Actually store the name-to-address mapping
 - Maintained by the corresponding administrative authority

Server hierarchy

- Each server stores a (small!) subset of the total DNS database
- An authoritative DNS server stores “**resource records**” for all DNS names in the domain that it has authority for
- Each server needs to know other servers responsible for other portions of the hierarchy
 - Every server knows the root
 - Root server knows about all top-level domains

DNS root

- Located in Virginia, USA
- How do we make the root scale?

13 DNS root servers



As of 2023-09-13T17:51:42Z, the root server system consists of 1751 instances operated by the 12 independent root server operators.

<https://root-servers.org/>

DNS records

- DNS servers store **resource records (RRs)**
 - RR is (name, value, type, TTL)
- Type = A: (→ **A**ddress)
 - name = hostname
 - value = IP address
- Type = NS: (→ **N**ame **S**erver)
 - name = domain
 - value = name of DNS server for domain

DNS records (cont'd)

- Type = CNAME: (→ Canonical Name)
 - name = alias name for some “canonical” (real) name
 - e.g., cse.umich.edu is really cse.eecs.umich.edu
 - value = canonical name
- Type = MX: (→ Mail eXchanger)
 - name = domain in email address
 - value = name(s) of mail server(s)

Inserting Resource Records into DNS

- Register foobar.com at registrar
 - Provide registrar with names and IP addresses of your authoritative name server(s)
 - Registrar inserts RR pairs into the .com TLD server:
 - (foobar.com, dns1.foobar.com, NS)
 - (dns1.foobar.com, 212.44.9.129, A)
- Store resource records in your server dns1.foobar.com
 - e.g., type A record for www.foobar.com
 - e.g., type MX record for foobar.com

Using DNS (Client/App View)

- Two components
 - Local DNS servers
 - Resolver software on hosts
- Local DNS server (“default name server”)
 - Clients configured with default server’s address OR learn it via a host configuration protocol (e.g., DHCP)

```
Wireless LAN adapter Wi-Fi:

Connection-specific DNS Suffix  . : adsroot.itcs.umich.edu
Description . . . . . : Intel(R) Wi-Fi 6E AX211 160MHz
Physical Address. . . . . : 00-A5-54-3B-A7-B3
DHCP Enabled. . . . . : Yes
Autoconfiguration Enabled . . . . : Yes
Link-local IPv6 Address . . . . . : fe80::b07e:59d8:3fed:5cb1%24(Preferred)
IPv4 Address. . . . . : [REDACTED] (Preferred)
Subnet Mask . . . . . : 255.255.0.0
Lease Obtained. . . . . : Wednesday, September 11, 2024 9:56:36 AM
Lease Expires . . . . . : Wednesday, September 11, 2024 1:29:06 PM
Default Gateway . . . . . : 35.3.0.1
DHCP Server . . . . . : 141.211.147.232
DHCPv6 IAID . . . . . : 167814484
DHCPv6 Client DUID. . . . . : 00-01-00-01-2B-34-E2-F8-00-0C-29-15-E0-FE
DNS Servers . . . . . : 10.10.10.10
                       : 10.10.5.5
NetBIOS over Tcpi. . . . . : Enabled
```

Using DNS (Client/App View)

- Two components
 - Local DNS servers
 - Resolver software on hostsadd
- Client application
 - Obtain DNS name (e.g., from URL)
 - Do `getnameinfo()` to trigger DNS request to its local DNS server

```
import socket  
s = socket.getaddrinfo("www.google.com",80)  
print(s)
```

```
[(<AddressFamily.AF_INET: 2>, 0, 0, '', ('142.250.190.132', 80))]
```

Dig (domain information groper)

`dig www.google.com`

```
; <<>> DiG 9.18.1-1ubuntu1.2-Ubuntu <<>> www.google.com
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 26427
;; flags: qr rd ad; QUERY: 1, ANSWER: 9, AUTHORITY: 0, ADDITIONAL: 0
;; WARNING: recursion requested but not available
```

```
;; QUESTION SECTION:
```


```
;www.google.com.                IN      A
```

```
;; ANSWER SECTION:
```


www.google.com.	300	IN	A	142.250.190.132
ns3.google.com.	0	IN	A	216.239.36.10
ns2.google.com.	0	IN	A	216.239.34.10
ns4.google.com.	0	IN	A	216.239.38.10
ns1.google.com.	0	IN	A	216.239.32.10

```
;; Query time: 9 msec
;; SERVER: 172.20.64.1#53(172.20.64.1) (UDP)
;; WHEN: Wed Sep 11 13:22:13 EDT 2024
;; MSG SIZE rcvd: 350
```



Name resolution




Name resolution




Name resolution




Name resolution



Name resolution: Recursive



Name resolution: Iterative



Two ways to resolve a name

- Recursive name resolution
 - Ask server to do it for you
- Iterative name resolution
 - Ask server who to ask next
- The iterative example we saw is a mix of both!

DNS protocol

- **Query** and **Reply** messages; both with the same message format
 - Header: identifier, flags, etc.
 - Plus resource records
- Client–server interaction on UDP Port 53
 - Spec supports TCP too, but not always implemented

Goals: Are we there yet?

- Uniqueness: No naming conflicts
- Scalable
- Distributed, autonomous administration
- Highly available?

Reliability

- **Replicated** DNS servers (primary/secondary)
 - Name service available if at least one replica is up
 - Queries can be load-balanced between replicas
- Usually, UDP used for queries
 - Reliability, if needed, must be implemented on UDP
- Try alternate servers on timeout
 - **Exponential backoff** when retrying same server
- Same identifier for all queries
 - Don't care which server responds

Goals: Are we there yet?

- Uniqueness: No naming conflicts
- Scalable
- Distributed, autonomous administration
- Highly available
- Fast lookups?

DNS caching

- Performing all these queries takes time
 - Up to 1-second latency before starting download
- Caching can greatly reduce overhead
 - The top-level servers very rarely change
 - Popular sites (e.g., www.google.com) visited often
 - Local DNS server often has the information cached
- How DNS caching works
 - DNS servers cache responses to queries
 - Responses include a “time to live” (TTL) field
 - Server deletes cached entry after TTL expires

TTL in dig output

```
dig www.google.com
```

```
; <<>> DiG 9.18.1-1ubuntu1.2-Ubuntu <<>> www.google.com
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 26427
;; flags: qr rd ad; QUERY: 1, ANSWER: 9, AUTHORITY: 0, ADDITIONAL: 0
;; WARNING: recursion requested but not available
```

```
;; QUESTION SECTION:
```

```
;www.google.com.                IN      A
```

```
;; ANSWER SECTION:
```

	TTL			
www.google.com.	300	IN	A	142.250.190.132
ns3.google.com.	0	IN	A	216.239.36.10
ns2.google.com.	0	IN	A	216.239.34.10
ns4.google.com.	0	IN	A	216.239.38.10
ns1.google.com.	0	IN	A	216.239.32.10

```
;; Query time: 9 msec
```

```
;; SERVER: 172.20.64.1#53(172.20.64.1) (UDP)
```

```
;; WHEN: Wed Sep 11 13:22:13 EDT 2024
```

```
;; MSG SIZE rcvd: 350
```

Negative caching

- Remember things that do not work
 - Misspellings like `www.google.comm`
 - These can take a long time to fail the first time
 - Good to remember that they do not work so the failure takes less time the next time around
- Negative caching is optional
 - Not widely implemented

Important properties of DNS

- Administrative delegation and hierarchy enables:
 - Easy unique naming
 - “Fate sharing” for network failures
 - Reasonable trust model
 - Caching increases scalability and performance

DNS provides indirection

- Addresses can change underneath
 - Move `www.cnn.com` to `4.125.91.21`
- Name could map to multiple IP addresses
 - Load-balancing ([CDN](#))
 - Reducing latency by picking nearby servers ([CDN](#))
 - Try out: `dig google.com` a few times
- Multiple names for the same address
 - E.g., many services (mail, www) on same machine
 - E.g., aliases like `www.cnn.com` and `cnn.com`

Summary

- CDNs improve web performance
 - Via replication and caching
 - Good server selection
- DNS allows us to go to webpages without having to memorize IP addresses
 - Allows a level of indirection that enables many functionalities including CDN server selection

Quiz 4

- <https://forms.gle/2rjqsa3SoGWdbdL58>

