



# **EECS 489**

## Computer Networks

---

Video streaming and Data Center Applications

# Agenda

- Video streaming
- Datacenter applications

# Ideas on video streaming

# How is video different?

- Often **too large** to send in one GET
- Doesn't even make sense even if it's possible
  - Users may skip forward!  $\Rightarrow$  save bandwidth wastage
  - Users' connection quality may change (e.g., switching from WiFi to LTE)  $\Rightarrow$  lower resolution to save bandwidth
- **Our focus is on stored video (i.e., not live)**

# Why video is important?

- Dominates the global Internet traffic landscape
  - About 60%, i.e., every 3 of 5 bytes in 2020!
  - About 65% in 2022!
    - Or 82% depending on the source
- Major sources
  - Netflix
  - YouTube
  - ...

# The video medium

- Video is a sequence of images/frames displayed at a constant rate (moving pictures)
- Digital image is an array of pixels, each pixel represented by bits
- Examples:
  - Single frame image encoding: 1024x1024 pixels, 24 bits/pixel  $\Rightarrow$  3 MB/image
  - Movies: 24 frames/sec  $\Rightarrow$  72 MB/sec
  - TV: 30 frames/sec  $\Rightarrow$  90 MB/sec

# The video medium (cont'd)

- Compression is key
  - Lots of algorithms to compress
- The same video can be (and typically is) compressed to multiple quality levels
  - E.g., 480p, 720p, 1080p, 4K
- Why multiple resolutions?
  - Adapt to conditions

# How do we serve video?

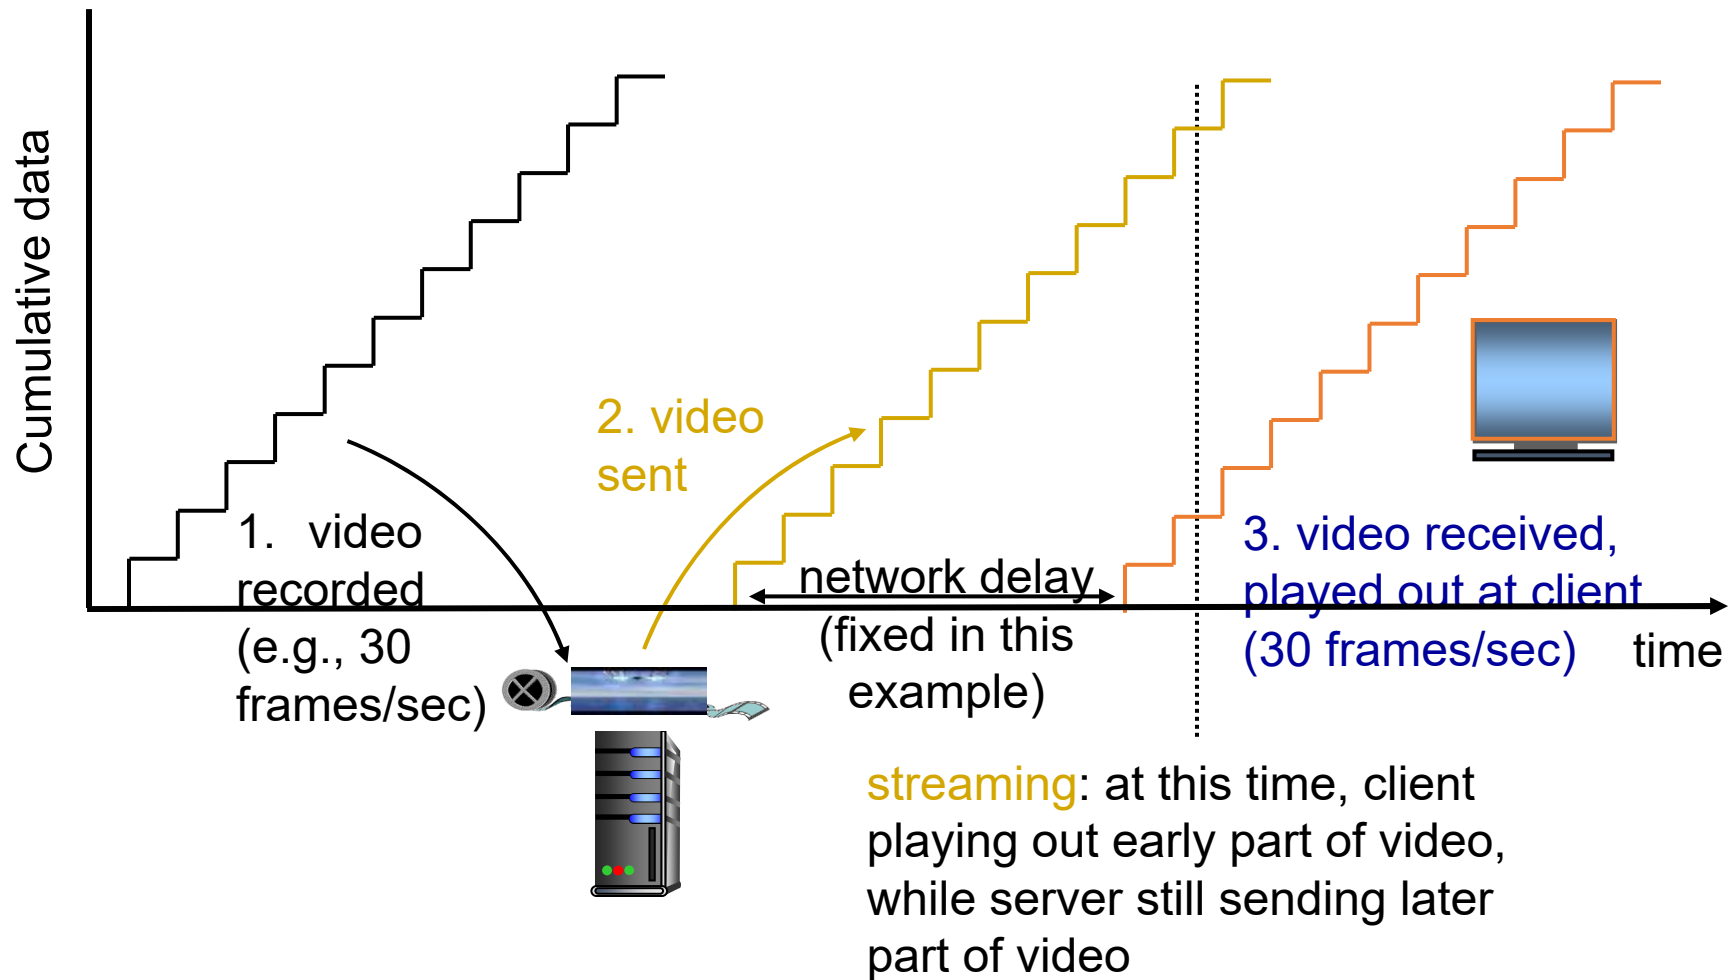
- It's in the name!
  - Video [streaming](#)



# HTTP streaming

- Video is stored at an HTTP server with a URL
- Clients send a GET request for the URL
- Server sends the video file as a stream
- Client first buffers for a while. Why?
  - To minimize interruptions later
- Once the buffer reaches a threshold
  - The video plays in the foreground
  - More frames are downloaded in the background

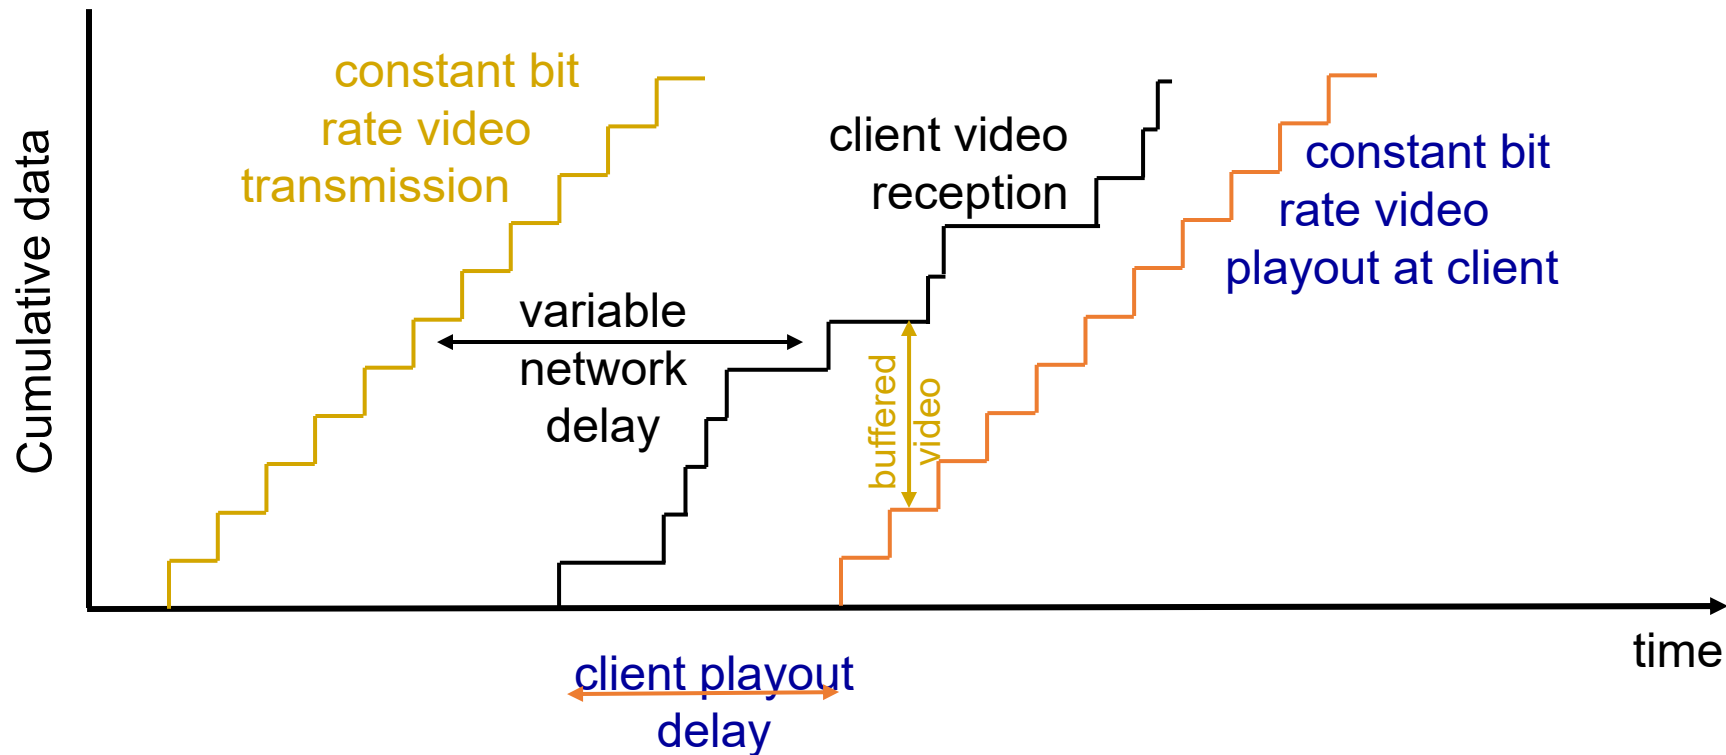
# HTTP streaming



# Challenges

- Absorb network delay variations
- Handle user interactions
  - Jump forward, fast-forward, rewind, pause
- Handle packet loss, retransmission etc.

# HTTP streaming: Revisited



# Issues with HTTP streaming


- Same bitrate for all clients
  - Clients can have very different network conditions
  - Clients network conditions can change over time
- Cannot dynamically adapt to conditions

# DASH : Dynamic Adaptive Streaming over HTTP

- Keep multiple resolutions of the same video
  - Stored in a manifest file in the HTTP server
- Client asks for the manifest file first to learn about the options
- Asks for chunks at a time and measures available bandwidth while they are downloaded
  - Low bandwidth  $\Rightarrow$  switch to lower bitrate
  - High bandwidth  $\Rightarrow$  switch to higher bitrate


# Cloud Systems

# Who's serving Web services?





# Who's serving Web services?




# Cloud datacenters run the world



Source: Google



# Cloud datacenters run the world





# Cloud datacenters run the world



Source: Facebook

# How big is a datacenter (DC)?

- 1M servers/site [Microsoft/Amazon/Google]
- > \$1B to build one site [Facebook]
- >\$20M/month/site operational costs [MS'09]
- Data center hardware spending grew to **\$177 billion** in 2017. [Gartner report]
- \$220 Billion in 2022
  
- But only  $O(10-100)$  sites

# Implications (1)

- Scale
  - Need scalable designs
  - Low-cost designs: e.g., use commodity technology
  - High utilization (efficiency): e.g., >60% avg. utilization
    - **Contrast**: avg. utilization on Internet links often ~30%
  - Tolerate frequent failure
    - Large number of (low cost) components
  - Automate

## Implications (2)

- Service model: clouds / multi-tenancy
  - Performance guarantees
  - Isolation guarantees
  - Portability


# Applications

- Common theme: **parallelism**
  - Applications decomposed into tasks
  - Running in parallel on different machines
- Two common paradigms
  - Partition-Aggregate
  - Map-Reduce



# Partition-Aggregate

eecs 489



# Partition-Aggregate

eeecs 489



GitHub - mosharaf/eeecs489: EECS 489: Computer Networks @ the ...

<https://github.com/mosharaf/eeecs489> ▼

**EECS 489**: Computer Networks (F'18) ... **EECS 489** takes a top-down approach to explore how networks operate and how network applications are written. ... Kurose and Ross, Computer Networking: A Top-Down Approach, 7th.

UM EECS 489: Computer Networks

[www.eecs.umich.edu/courses/eeecs489/](http://www.eecs.umich.edu/courses/eeecs489/) ▼

Lecture: MWF 9:30 - 10:30 in 1500 **EECS**. Discussion/Lab: W 12:30 - 1:30 in 2166 DOW or W 4:30 - 5:30 in 1014 DOW. The discussion sessions will mostly be ...

EECS 489

[www.eecs.umich.edu/courses/eeecs489/f99/](http://www.eecs.umich.edu/courses/eeecs489/f99/) ▼

News group umich.**eeecs**.class.489. Everything posted here will be automatically forwarded to the **eeecs489staff@eeecs.umich.edu** mailing list hourly.

EECS 489 - EECS @ Michigan - University of Michigan

<https://www.eecs.umich.edu/eeecs/academics/courses/eeecs-489.html> ▼

Course Homepage: <http://www.eecs.umich.edu/courses/eeecs489/w10/>. Coverage We study how networks operate and how network applications are written.

Revamping EECS 489: A Retrospective | Mosharaf Chowdhury

<https://www.mosharaf.com/blog/2017/05/07/revamping-eeecs-489-a-retrospective/> ▼


May 7, 2017 - A couple of weeks ago, we wrapped up the Spring 2017 offering of the **EECS 489**:



# End-to-end response time

- **Less than 200 milliseconds** between receiving user query in the browser and displaying the results
  - $RTT = O(10)$  to 100 milliseconds
  - What remains?
- Next time, when the page is not loading fast enough, think about the poor servers working for you 😊

# Map-Reduce




The most popular software that follows this paradigm is [Apache Spark](#)

# Announcements


- Midterm is Wed October 16 2024
  - 6PM – 8PM (in-person)
  - Room assignments to be announced later
  - Please email me if you have a conflict with another exam
- Assignment 1 due Wednesday, 11:59 PM.

5-minute break!

# Datacenter networks




## Datacenter networks (Cont.)





# Datacenter traffic



# East-West traffic

- Traffic **between servers** in the datacenter
- Communication within “big data” computations
- Traffic may shift on small timescales (< minutes)


# Datacenter traffic characteristics

- Two key characteristics
  - Most flows are small
  - Most bytes come from large flows
- Applications want
  - High bandwidth (large flows)
  - Low latency (small flows)


# High bandwidth

- Ideal: Each server can talk to any other server at its full access link rate
- Conceptually: Datacenter network as one giant switch


# Datacenter network as one giant switch



# Datacenter network as one giant switch



# Datacenter network as one giant switch



# High bandwidth

- Ideal: Each server can talk to any other server at its full access link rate
- Conceptually: Datacenter network as one giant switch
  - Would require a 10 Pbits/sec switch! ( $10^{15}$  bps)
    - 1M ports (one port/server)
    - 10Gbps per port
- **Practical approach:** build a network of switches (“fabric”) with high “bisection bandwidth”
  - Each switch has practical #ports and link speeds




# Bisection bandwidth

- Partition a network into two equal parts
- Minimum bandwidth between the partitions is the bisection bandwidth
- **Full bisection bandwidth**: bisection bandwidth in an  $N$  node network is  $N/2$  times the bandwidth of a single link
  - Nodes of any two halves can communicate at full speed with each other

# Achieving full bisection bandwidth

- Scale up
  - Make links fatter toward the core of the network
- Problem: Scaling up a traditional tree topology is expensive!
  - Requires non-commodity / impractical / link and switch components
- Solutions?
  - Over-subscribe (i.e., provision less than full BW)
  - Better topologies


# Oversubscription




# Oversubscription

- Not enough bandwidth
  - **Oversubscription**: Less bandwidth in the ToR-Agg links than all the servers bandwidth in the rack
  - **Oversubscription ratio**: Ratio between bandwidth underneath and bandwidth above
- Not enough paths between server pairs
  - Load balancing issues
  - Failure recovery issues

# Better topologies




# Better topologies




# Clos topology

- Multi-stage network
- $k$  pods, where each pod has two layers of  $k/2$  switches (k-ports)
  - $k/2$  ports up and  $k/2$  down
- All links have the same b/w
- At most  $k^3/4$  machines
- Example
  - $k = 4$
  - 16 machines
- For  $k=48$ , 27648 machines



# Challenges in scale-out designs?

- Topology offers high bisection bandwidth
- All other system components must be able to exploit this available capacity
  - Routing must use all paths
  - Transport protocol must fill all pipes (fast)





# Summary

- Video streaming
  - Too large to send, so stream it
  - Dynamically adapt to the network and users
- Cloud systems
  - Forms the backend of modern web services
  - Runs in datacenters where all the processing happens

# Quiz

- <https://forms.gle/UY6ZpYBgtDEFuAHE9>

