



EECS 489

Computer Networks

Transport Control Protocol - TCP

TCP: Transmission Control Protocol

Build the TCP header

Source port	Destination port
Sequence number	
Acknowledgment	

Checksum

Data

What does TCP do?

- Most of what we've seen
 - Checksum
 - Sequence numbers are byte offsets
 - Receiver sends cumulative acknowledgements (like GBN)
 - Receivers **can buffer out-of-sequence packets** (like SR)

What does TCP introduce?

- Most of what we've seen
 - Checksum
 - Sequence numbers are byte offsets
 - Receiver sends cumulative acknowledgements (like GBN)
 - Receivers buffer out-of-sequence packets (like SR)
- Introduces fast retransmit: duplicate ACKs trigger early retransmission
- Sender maintains a **single retransmission timer** (like GBN) and retransmits on timeout

Review from Discussion 1:

- Send message

```
do {  
    ssize_t n = send(sockfd, message + sent, message_len - sent, 0);  
    sent += n;  
} while (sent < message_len);
```

- Receive message

```
do {  
    // Receive as many additional bytes as we can in one call to recv()  
    // (while not exceeding MAX_MESSAGE_SIZE bytes in total).  
    rval = recv(connectionfd, msg + recvd, MAX_MESSAGE_SIZE - recvd, 0);  
    recvd += rval;  
} while (rval > 0); // recv() returns 0 when client closes
```

HTTP exchange example

- Client connects to a webserver and generates a GET request

GET / HTTP/1.1

Accept:

text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image/webp,image/apng,*/*;q=0.8,application/signed-exchange;v=b3;q=0.7

Accept-Encoding: gzip, deflate

Accept-Language: en-US,en;q=0.9

Cache-Control: no-cache

Connection: keep-alive

Host: www.eecs489.org

Pragma: no-cache

Upgrade-Insecure-Requests: 1

User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64)

AppleWebKit/537.36 (KHTML, like Gecko) Chrome/116.0.0.0 Safari/537.36

HTTP exchange example

- Server sends a response

HTTP/1.1 200 OK

Connection: keep-alive

Content-Length: 5000


Server: GitHub.com

Content-Type: text/html; charset=utf-8

Last-Modified: Wed, 20 Sep 2024 17:31:52 GMT

... (Total of 5,000 bytes)

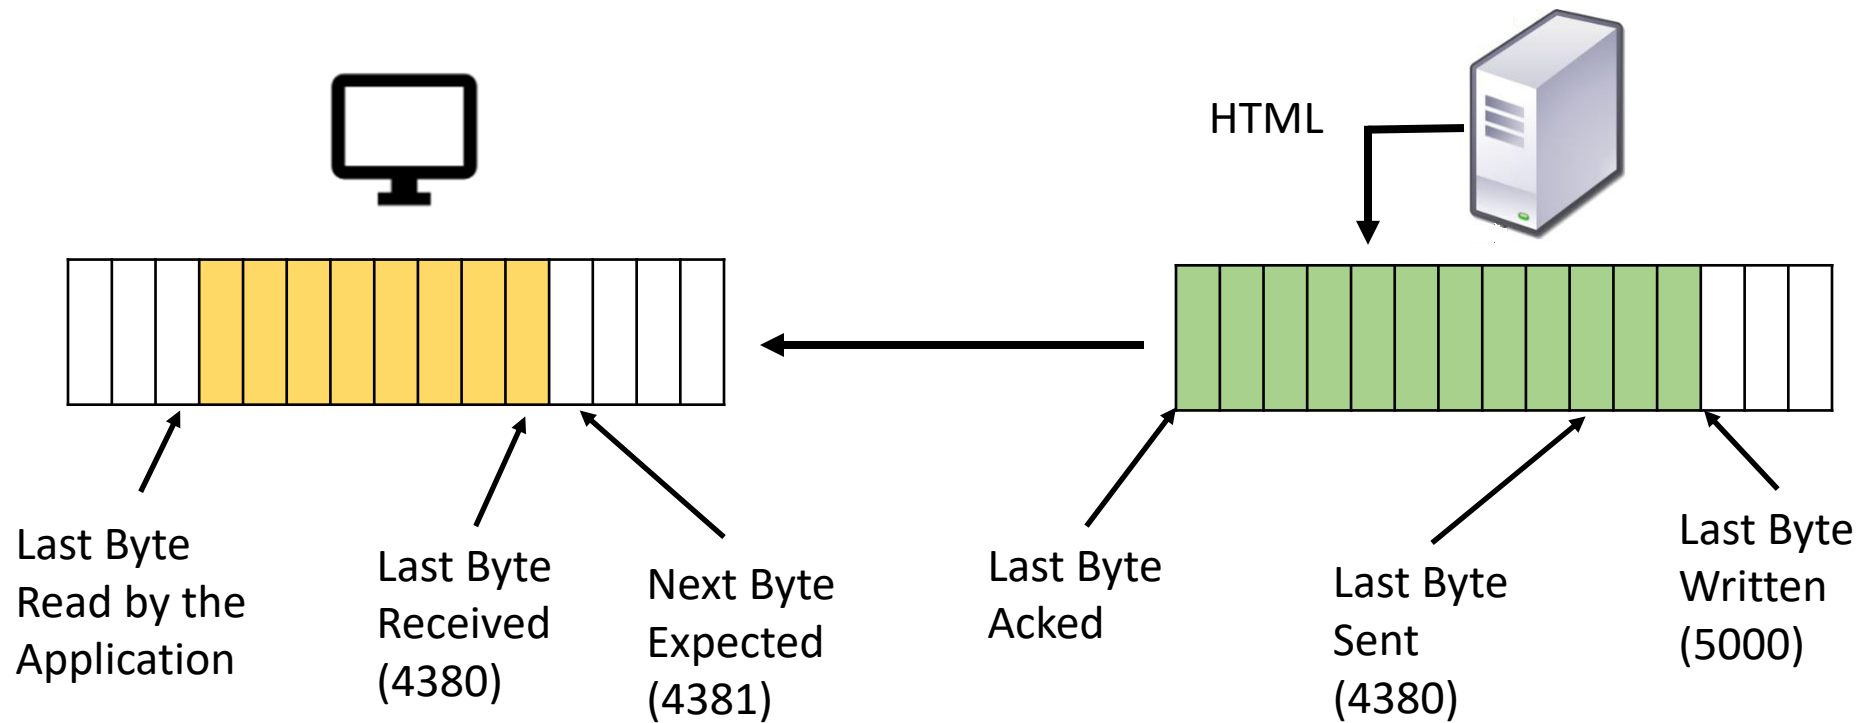
Now the server is going to send the webpage



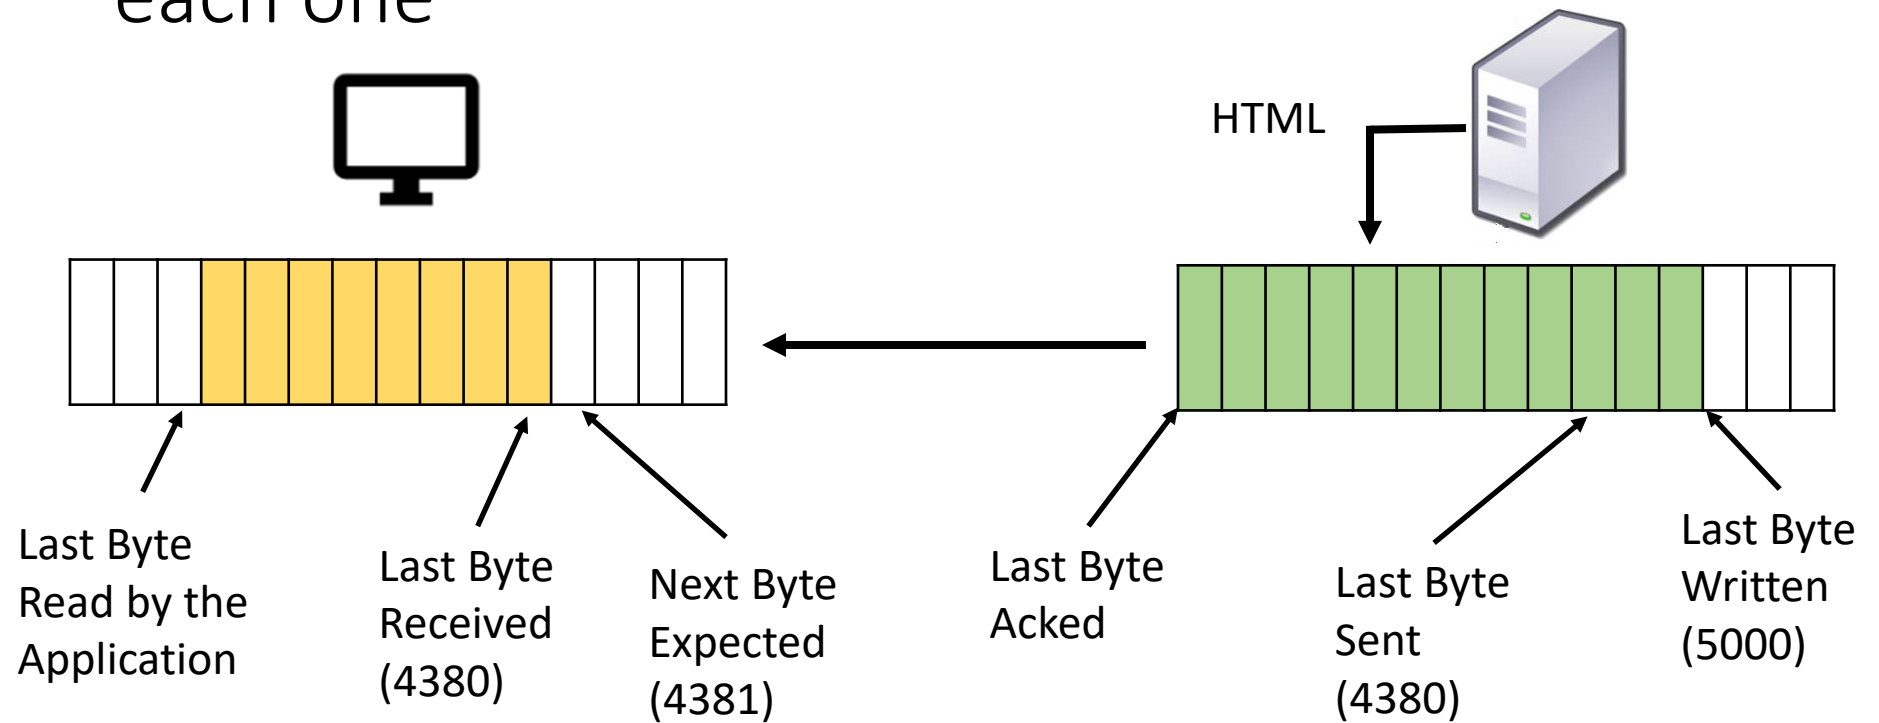
How much data can the server send

- Let's say for each MTU (Maximum Transmission Unit) is 1500 Bytes
- IP Header and TCP headers are 20 bytes each.
- Each TCP Packet can hold $1500 - 20 - 20 = 1460$ Bytes
- So TCP can transmit 1460 bytes of data in each packet
- Server's Sliding window is 5000 Bytes

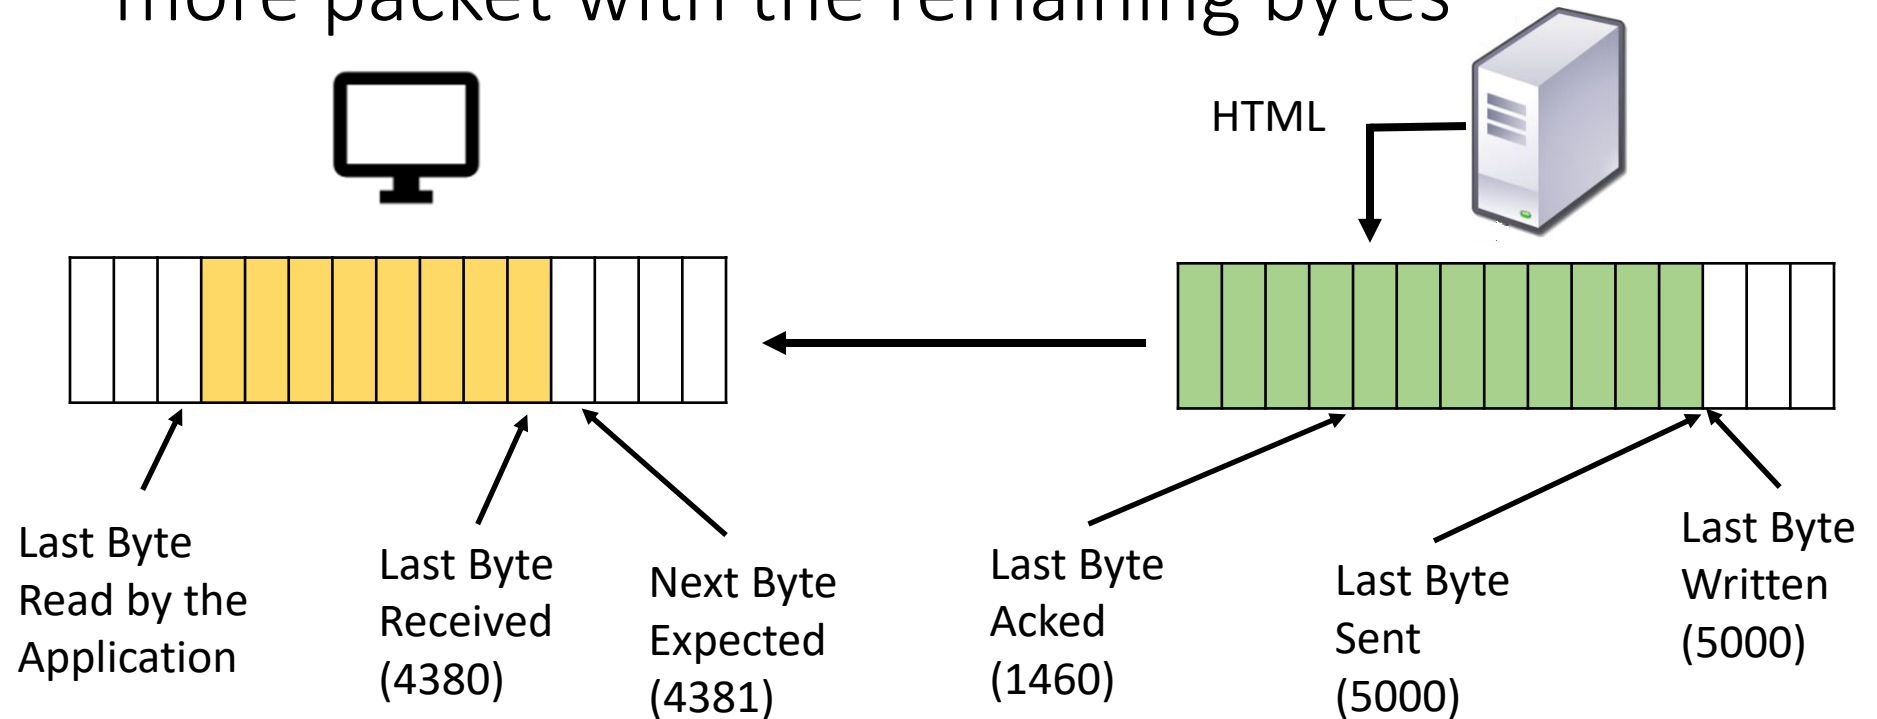
Server sends 3 packets, each with 1460 bytes



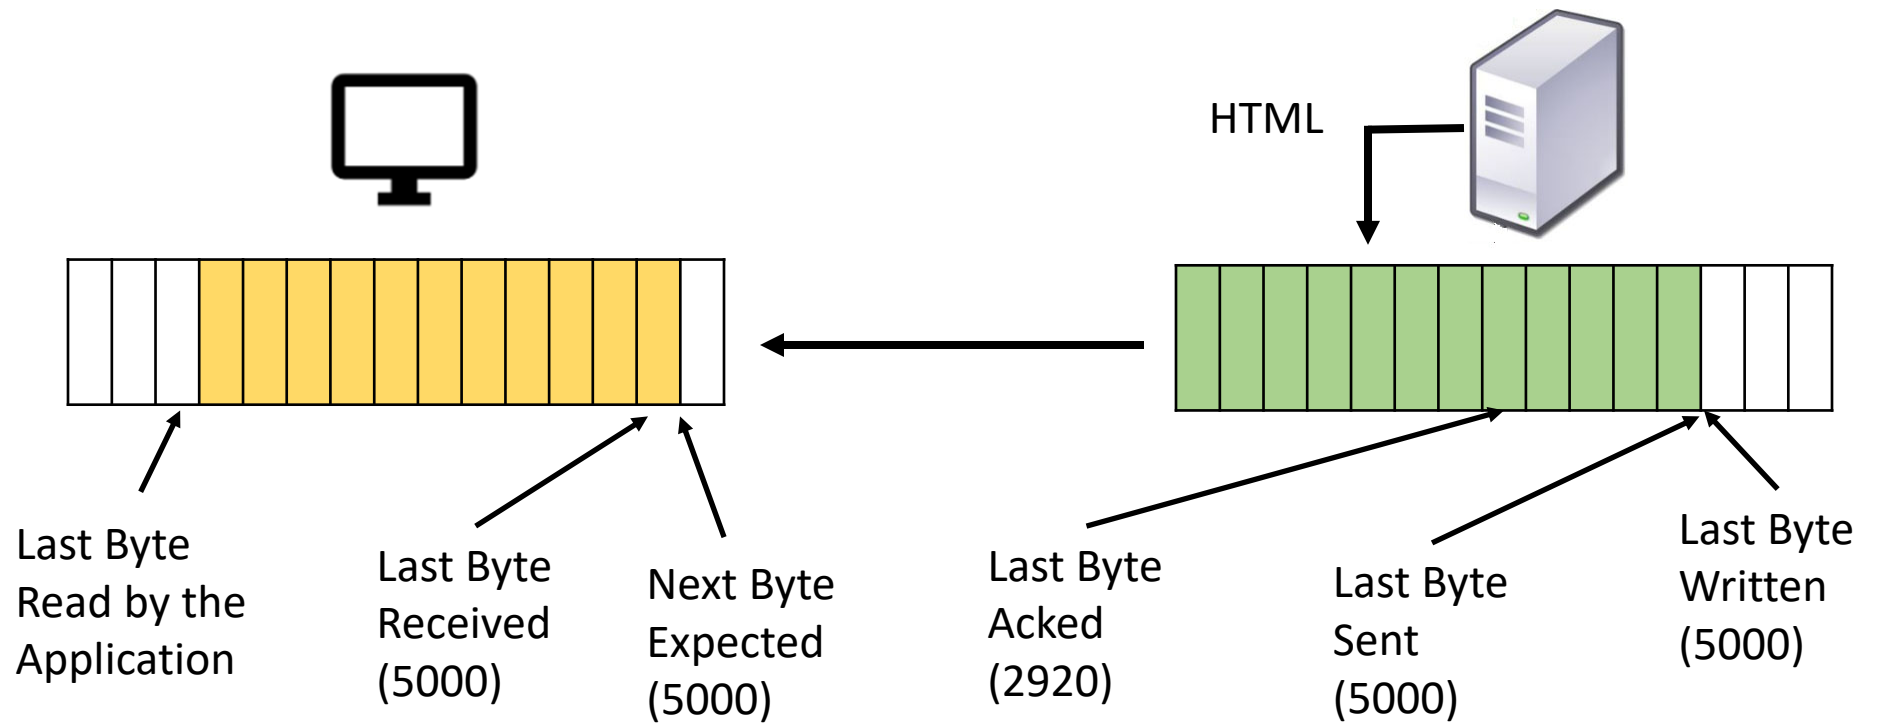
Host receives the three packets and ACKS each one



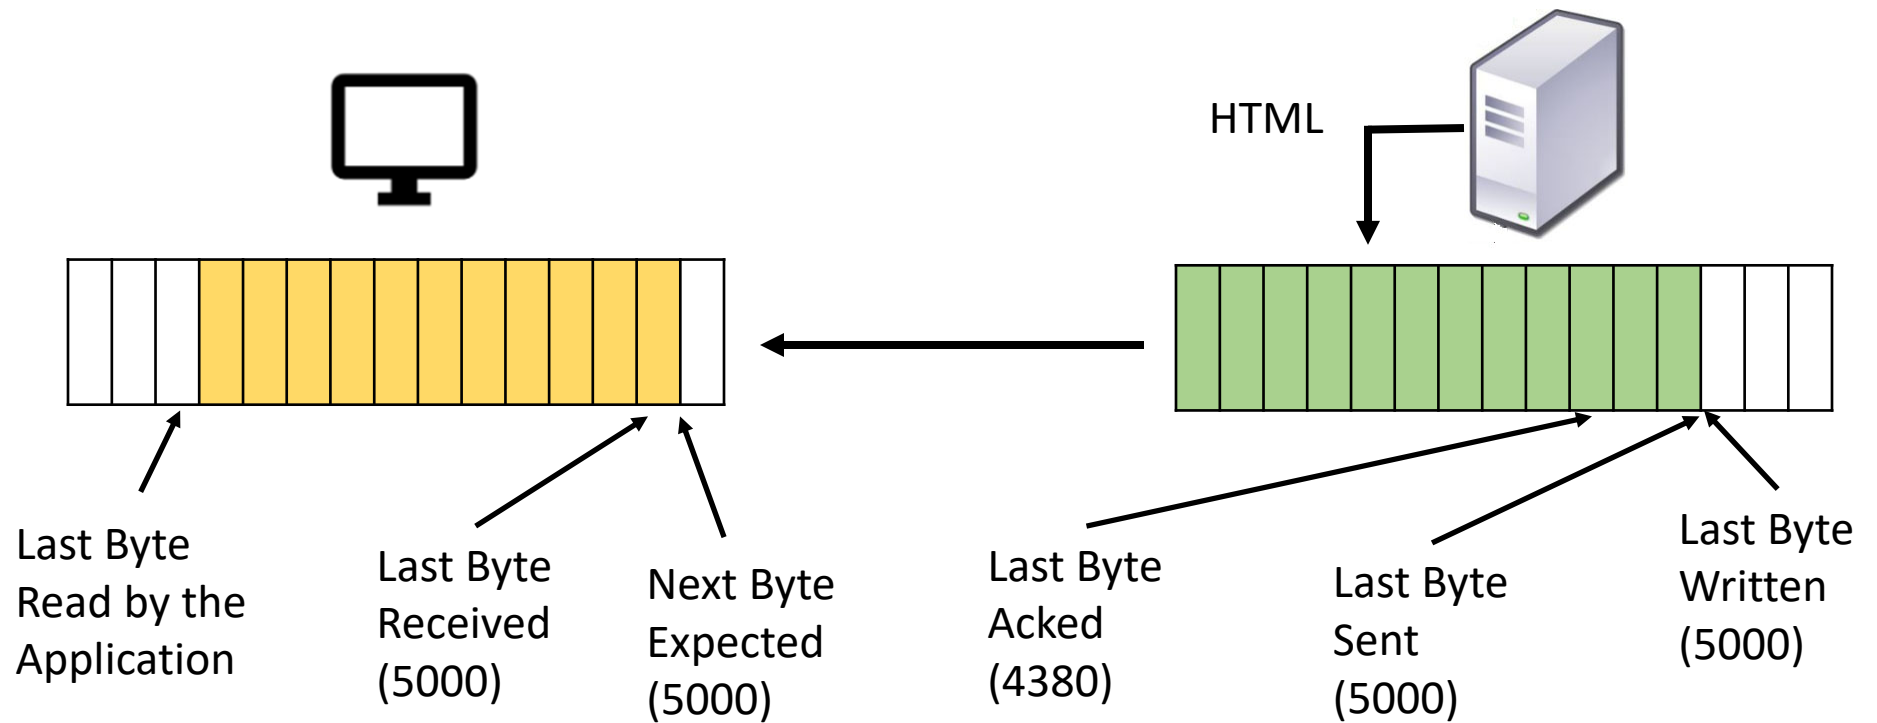
Server after receiving first ACK – send one more packet with the remaining bytes



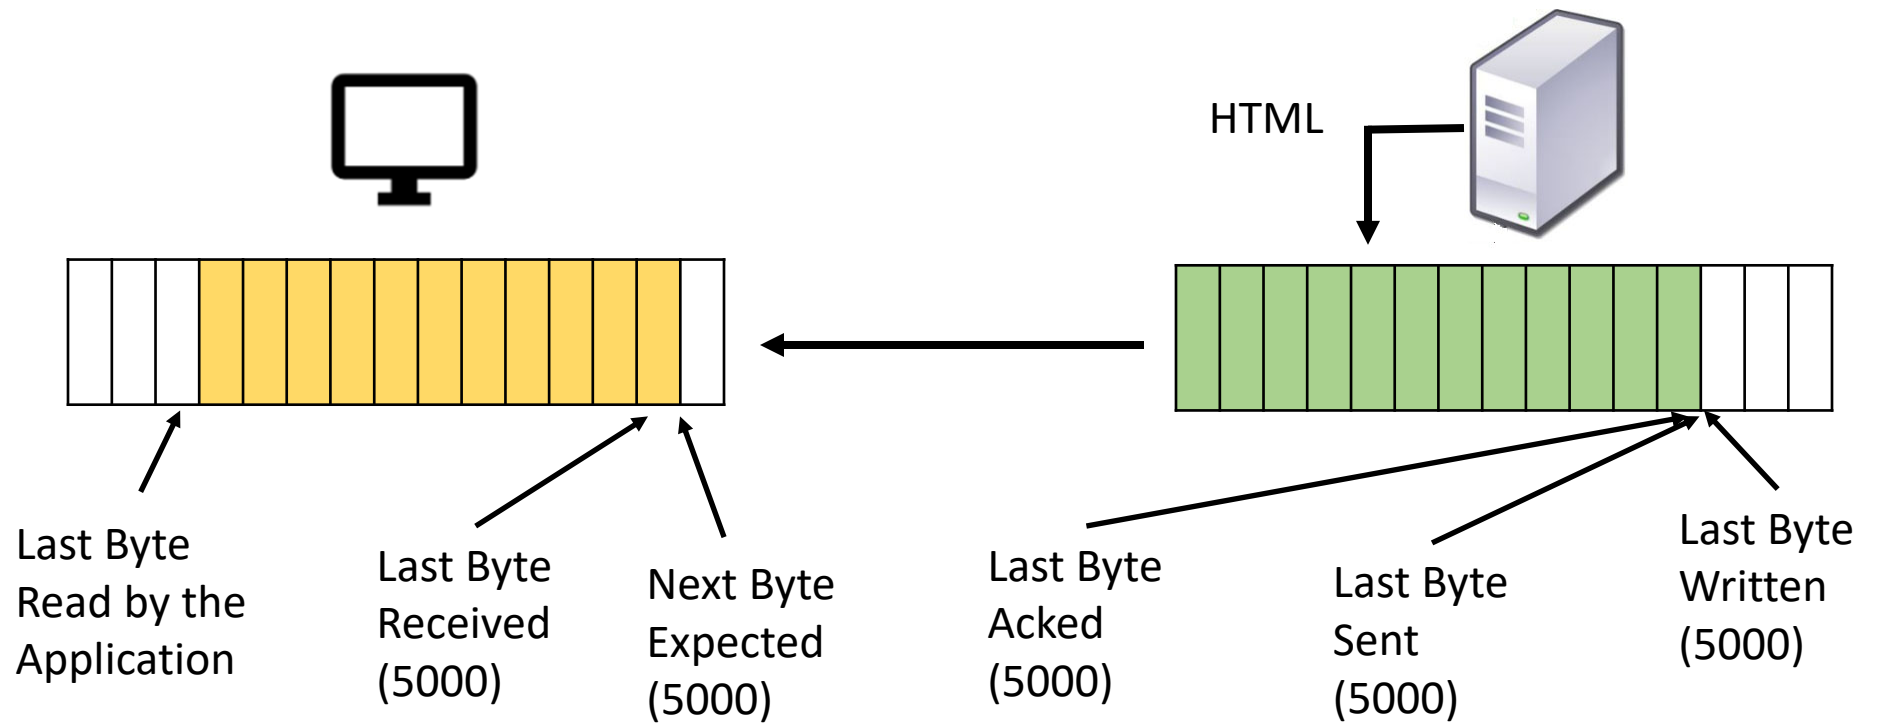
Server after receiving second ACK



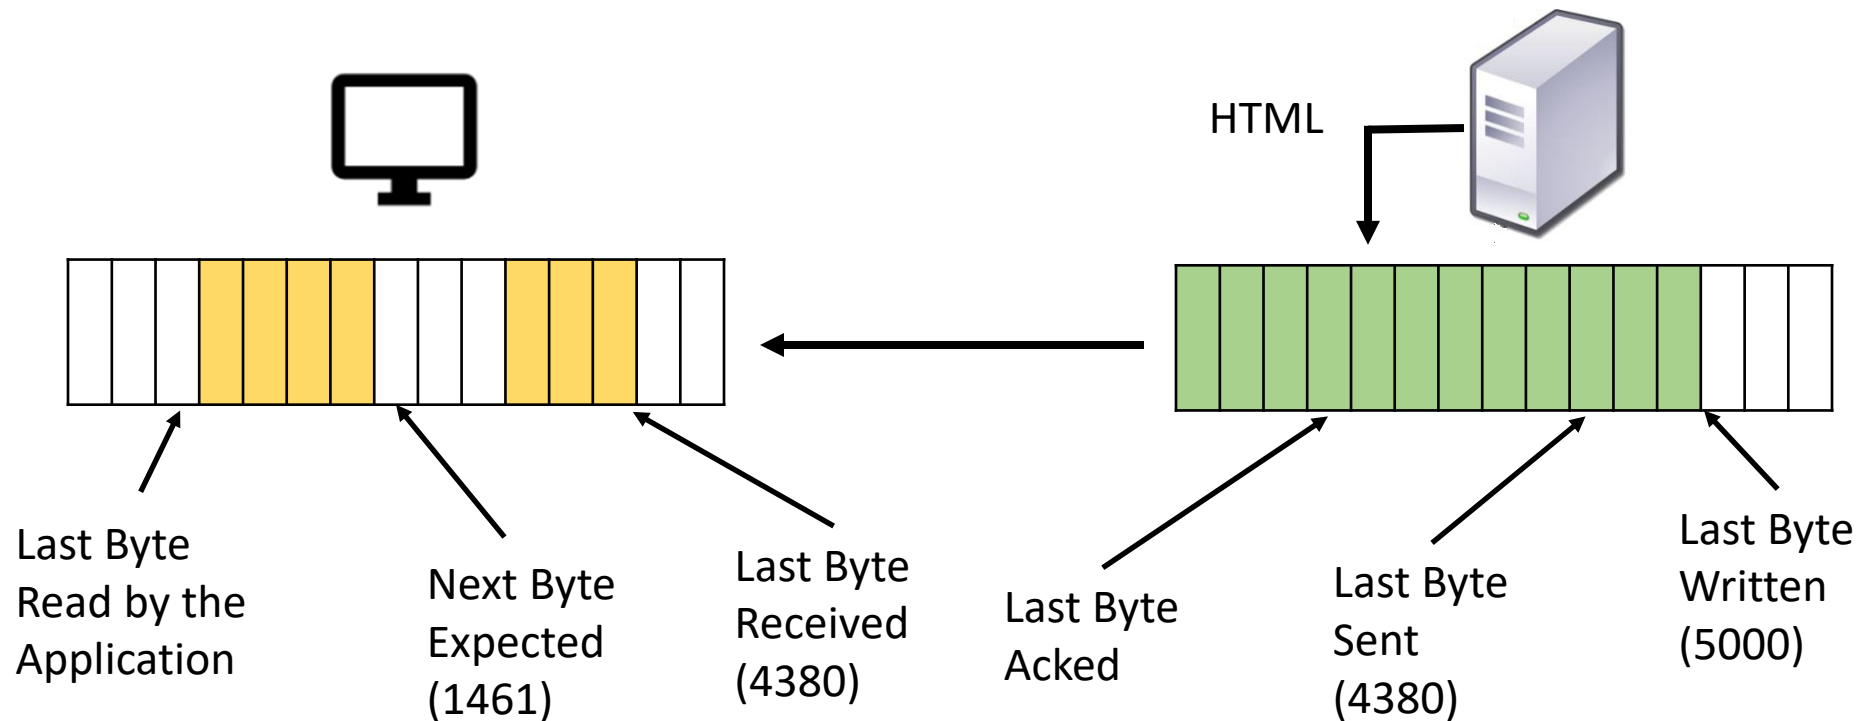
Server after receiving third ACK



Server after receiving fourth ACK



Out of Order Delivery/Missing packet




- Host Acks with 1461 on Receiving Bytes 2920 – 4380
- Server retransmits Bytes 1461 – 2919 on receiving the above ACK or on a timeout


Retransmission timeout

- If the sender hasn't received an ACK by timeout, **retransmit the first packet** in the window
- How do we pick a timeout value?

Timing illustration



Timeout too long → inefficient



Timeout too short →
duplicate packets


Retransmission timeout

- If the sender hasn't received an ACK by timeout, retransmit the first packet in the window
- How to set timeout?
 - Too long: connection has low throughput
 - Too short: retransmit packet that was just delayed
- Solution: **make timeout proportional to RTT**
 - But how do we measure RTT?

RTT estimation


- Exponential weighted average of RTT samples

$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$



Problem: Ambiguous measurements

- How do we differentiate between the real ACK, and ACK of the retransmitted packet?



Karn/Partridge algorithm

- Don't use SampleRTT from retransmissions
 - Once retransmitted, ignore that segment in the future
- Computes EstimatedRTT using $\alpha = 0.125$
- Timeout value (RTO) = $2 \times \text{EstimatedRTT}$
 - Employs exponential backoff
 - Every time RTO timer expires, set $\text{RTO} \leftarrow 2 \cdot \text{RTO}$
(Up to maximum ≥ 60 sec)
 - Every time new measurement comes in (= successful original transmission), collapse RTO back to $2 \times \text{EstimatedRTT}$
- Sensitive to RTT variations

Jacobson/Karels algorithm

- **Problem:** need to better capture variability in RTT
 - Directly measure deviation

$$\text{Difference} = \text{SampleRTT} - \text{EstimatedRTT}$$

$$\text{EstimatedRTT} = \text{EstimatedRTT} + (\delta \times \text{Difference})$$

$$\text{Deviation} = \text{Deviation} + \delta(|\text{Difference}| - \text{Deviation})$$

$$\text{TimeOut} = \mu \times \text{EstimatedRTT} + \phi \times \text{Deviation}$$

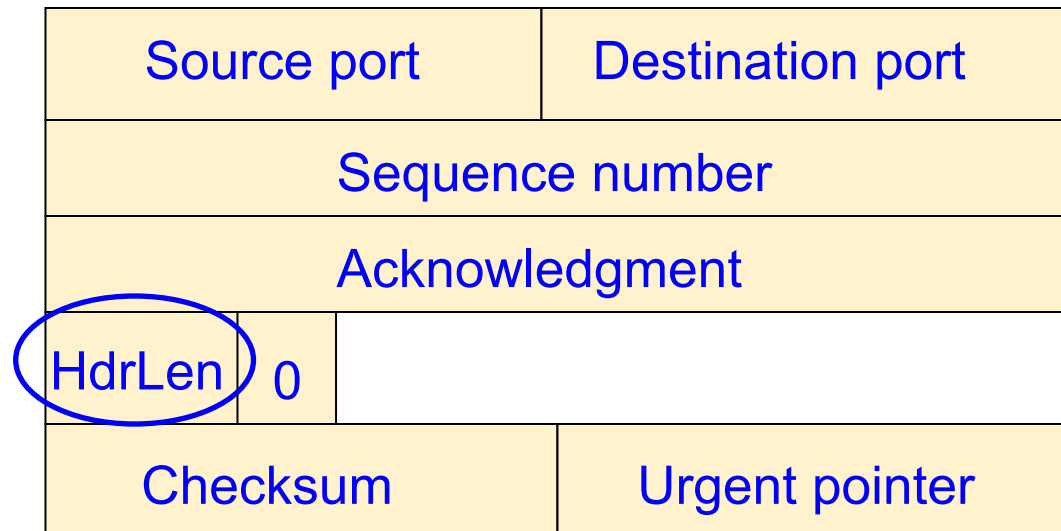
μ is typically set to 1 and ϕ is set to 4*

- **RTO = EstimatedRTT + 4 x DeviationRTT**

*Peterson and Davie

Build the TCP header

Number of 4-byte
words in the header;
5: No options




TCP Connection Establishment

Initial Sequence Number (ISN)

- Sequence number for the very first byte
- Why not just use ISN = 0?
 - Practical issue
 - IP addresses and port #s uniquely identify a connection
 - Eventually, though, these port #s do get used again; small chance an old packet is still in flight
 - Also, others might try to spoof your connection
 - Why does using ISN help?
- Hosts exchange ISNs when establishing connection

Establishing a TCP connection

- **Three-way handshake** to establish connection
 - Host A sends a SYN (open; “synchronize sequence numbers”) to host B
 - Host B returns a SYN acknowledgment (SYN ACK)
 - Host A sends an ACK to acknowledge the SYN ACK



Build the TCP header

Flags:

SYN

ACK

FIN

RST

PSH

URG

Source port		Destination port	
Sequence number			
Acknowledgment			
HdrLen	0	Flags	
Checksum		Urgent pointer	

Data

Step 1: A's initial SYN packet

A tells B to open
a connection

A's port			B's port		
A's Initial Sequence Number					
N/A					
5	0	SYN			
Checksum			Urgent pointer		

Step 1: B's SYN-ACK packet

B tells it accepts
and is ready to
accept next
packet

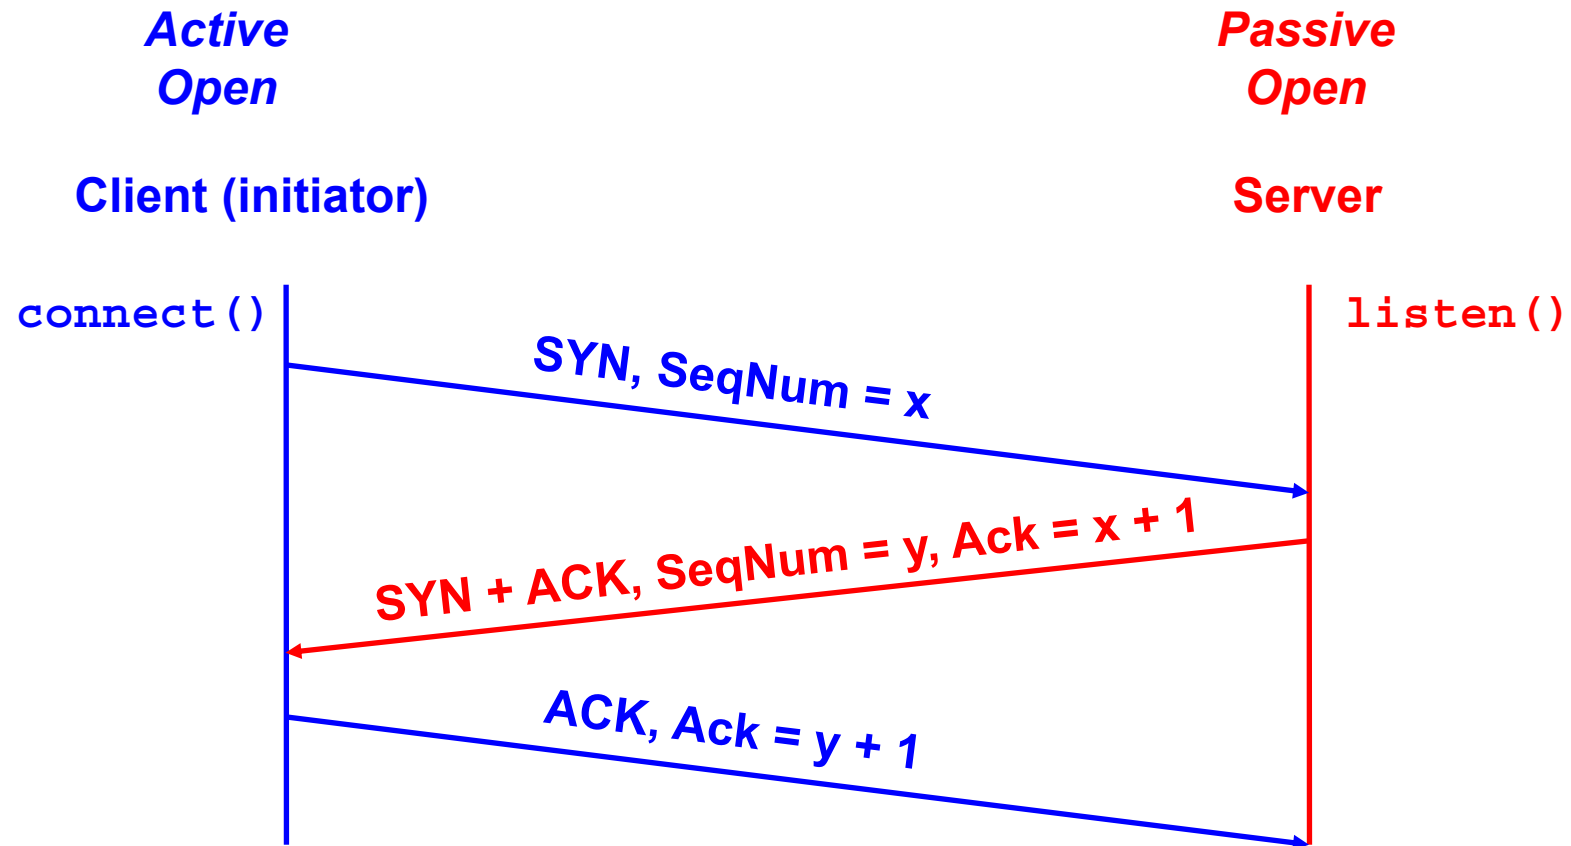
B's port		A's port	
B's Initial Sequence Number			
ACK=A's ISN+1			
5	0	SYN ACK	
Checksum		Urgent pointer	

Step 1: A's ACK to SYN-ACK

A tells B to open
a connection

A's port		B's port	
A's Initial Sequence Number + 1			
ACK=B's ISN+1			
5	0	ACK	
Checksum		Urgent pointer	

TCP's 3-Way handshaking



What if the SYN Packet Gets Lost?


- Suppose the SYN packet gets lost
 - Packet dropped by the network or server is busy
- Eventually, no SYN-ACK arrives
 - Sender retransmits the SYN on timeout
- How should the TCP sender set the timer?
 - Sender has no idea how far away the receiver is
 - Hard to guess a reasonable length of time to wait
 - SHOULD (RFCs 1122 & 2988) use default of 3 seconds
 - Some implementations instead use 6 seconds

SYN loss and web downloads


- User clicks on a hypertext link
 - Browser creates a socket and does a “connect”
 - The “connect” triggers the OS to transmit a SYN
- If the SYN is lost...
 - 3-6 seconds of delay: can be very long
 - User may become impatient and can retry
- User triggers an “abort” of the “connect”
 - Browser creates a new socket and another “connect”
 - Can be effective in some cases

TCP connection teardown

Normal termination, one side at a time




Normal termination, both together



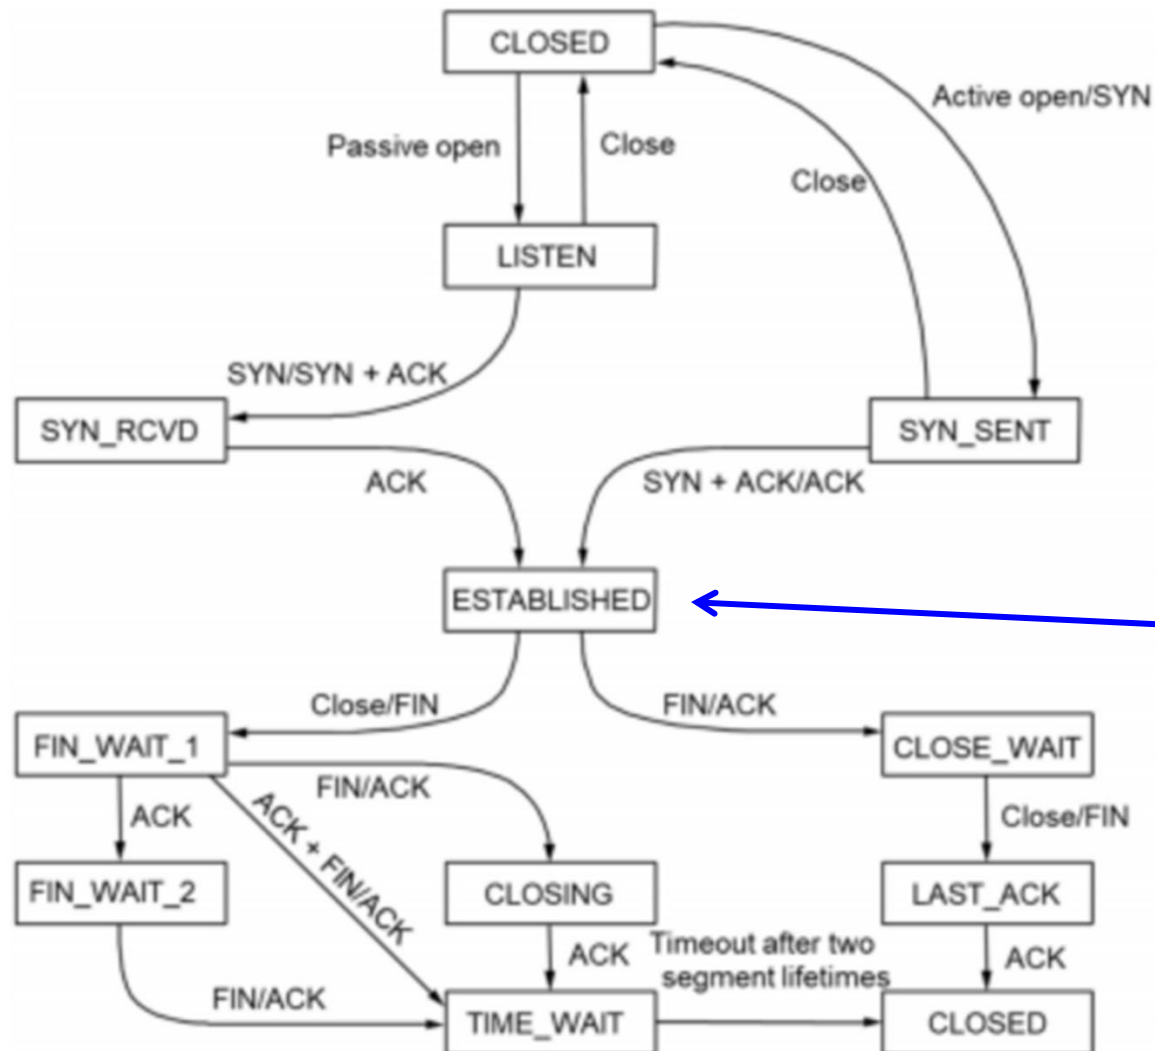
- Same as before, but B sets FIN with their ack of A's FIN

Abrupt termination



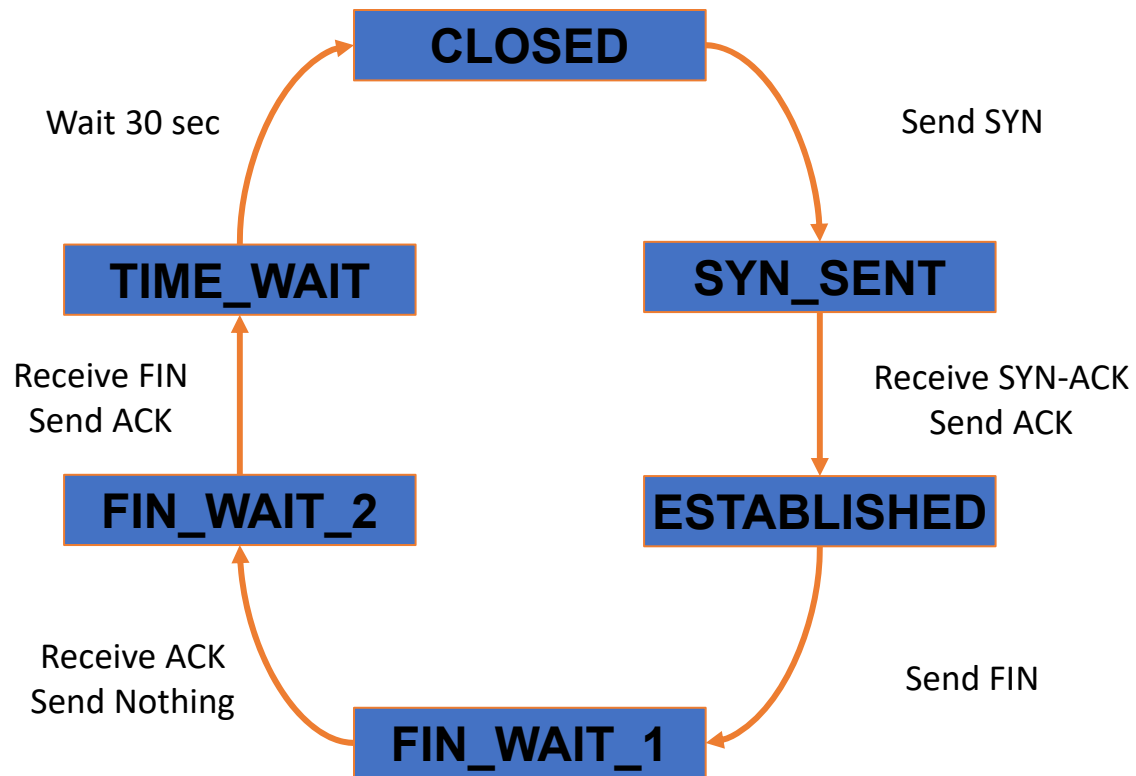
- A sends a RESET (RST) to B
 - E.g., because application process on A crashed
- That's it
 - B does not ack the RST
 - Thus, RST is not delivered reliably, and any data in flight is lost
 - But: if B sends anything more, will elicit another RST

TCP State Transition



Data, ACK exchanges are in here

TCP client lifecycle



HTTP exchange example

- Client connects to a webserver and generates a GET request

GET / HTTP/1.1

Accept:

text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image/webp,image/apng,*/*;q=0.8,application/signed-exchange;v=b3;q=0.7

Accept-Encoding: gzip, deflate

Accept-Language: en-US,en;q=0.9

Cache-Control: no-cache

Connection: keep-alive

Host: www.eecs489.org

Pragma: no-cache

Upgrade-Insecure-Requests: 1

User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64)

AppleWebKit/537.36 (KHTML, like Gecko) Chrome/116.0.0.0 Safari/537.36

HTTP exchange example

- Server sends a response

HTTP/1.1 200 OK

Connection: keep-alive

Content-Length: 10257

Server: GitHub.com

Content-Type: text/html; charset=utf-8

Last-Modified: Wed, 20 Sep 2023 17:31:52 GMT

... (Total of 53,568 bytes)

TCP Message Header

U	A	P	R	S	F
---	---	---	---	---	---



Source Port							Destination Port						
Sequence Number													
Acknowledgement													
Length	Reserved		U	A	P	R	S	F					
Checksum								Urgent Pointer					
Options (0-40 Bytes)													
Data (Optional)													

U	A	P	R	S	F
---	---	---	---	---	---

Application says connect to server

ISN: 4076393977



SYN



Randomly
Picked by
the OS for
Host

63220

80

Randomly
Picked by
TCP

4076393977

0

5

Reserved

000010

Checksum


Urgent Pointer

Options (0-40 Bytes)

Data (Optional)

U	A	P	R	S	F
---	---	---	---	---	---

Application says connect to server



80				63220			
2120529931							
4076393978							
ed		0	1	0	0	1	0
Checksum				Urgent Pointer			
Options (0-40 Bytes)							
Data (Optional)							

I have received all packets before this Seq Num


TCP on the server picks

I have received all packets before this Seq Num

TCP on the server picks an ISN

U	A	P	R	S	F
---	---	---	---	---	---


Application says connect to server



63220							80						
4076393978													
2120529932													
5	Reserved		0	1	0	0	0	0					
Checksum								Urgent Pointer					
Options (0-40 Bytes)													
Data (Optional)													

U	A	P	R	S	F
---	---	---	---	---	---


Application Sends data (451 Bytes)



62330										80																					
4076393979																															
2120529932																															
5	Reserved				0	0	0	0	0	0																					
Checksum										Urgent Pointer																					
Options (0-40 Bytes)																															
G	E	T		i	n	d	e	x	.	h	t	m	l		H	T	T	P	/	1	.	1		A	C	C	E	P	T	.	.

U	A	P	R	S	F
---	---	---	---	---	---

Application says connect to server



80								6230							
2120529932															
4076394493															
5		Reserved		0	1	0	0	0	0	Urgent Pointer					
Checksum															
Options (0-40 Bytes)															
Data (Optional)															

4076393979
+ 514

Bonus Quiz 7

- <https://forms.gle/iroh9cEyjUxi1L238>

