



# **EECS 489**

## **Computer Networks**

---

Routing Fundamentals

# Agenda

- Finishing up from last lecture
- Routing fundamentals

# Router-assisted congestion control


- Three tasks for congestion control
  - Isolation/fairness
  - Adjustment
  - Detecting congestion

# Fairness: General approach

- Routers classify packets into “flows”
  - Let’s assume flows are TCP connections
- Each flow has its own FIFO queue in router
- Router services flows in a fair fashion
  - When line becomes free, take packet from next flow in a fair order
- What does “fair” mean exactly?


# Max-Min fairness

- Given set of bandwidth demands  $r_i$  and total bandwidth  $C$ , max-min bandwidth allocations are:
  - $a_i = \min(f, r_i)$
  - where  $f$  is the unique value such that  $\text{Sum}(a_i) = C$



# Example

- $C = 10; r_1 = 8, r_2 = 6, r_3 = 2; N = 3$
- $C/3 = 3.33 \rightarrow$ 
  - $r_3$  needs only 2
    - Can service all of  $r_3$
    - Remove  $r_3$  from the accounting:  $C = C - r_3 = 8; N = 2$
- $C/2 = 4 \rightarrow$ 
  - Can't service all of  $r_1$  or  $r_2$
  - So hold them to the remaining fair share:  $f = 4$



$f = 4:$   
 $\min(8, 4) = 4$   
 $\min(6, 4) = 4$   
 $\min(2, 4) = 2$

# Max-Min fairness

- Given set of bandwidth demands  $r_i$  and total bandwidth  $C$ , max-min bandwidth allocations are:
  - $a_i = \min(f, r_i)$
  - where  $f$  is the unique value such that  $\text{Sum}(a_i) = C$
- If you don't get full demand, no one gets more than you
- This is what round-robin service gives if all packets are the same size

# How do we deal with packets of different sizes?





- Mental model: Bit-by-bit round robin (“fluid flow”)
- Can you do this in practice?
  - No, packets cannot be preempted
- But we can approximate it
  - This is what “fair queuing” routers do



# Fair Queuing (FQ)

- For each packet, compute the time at which the last bit of a packet would have left the router if flows are served bit-by-bit
- Then serve packets in the increasing order of their deadlines

# Example



# Fair Queuing (FQ)

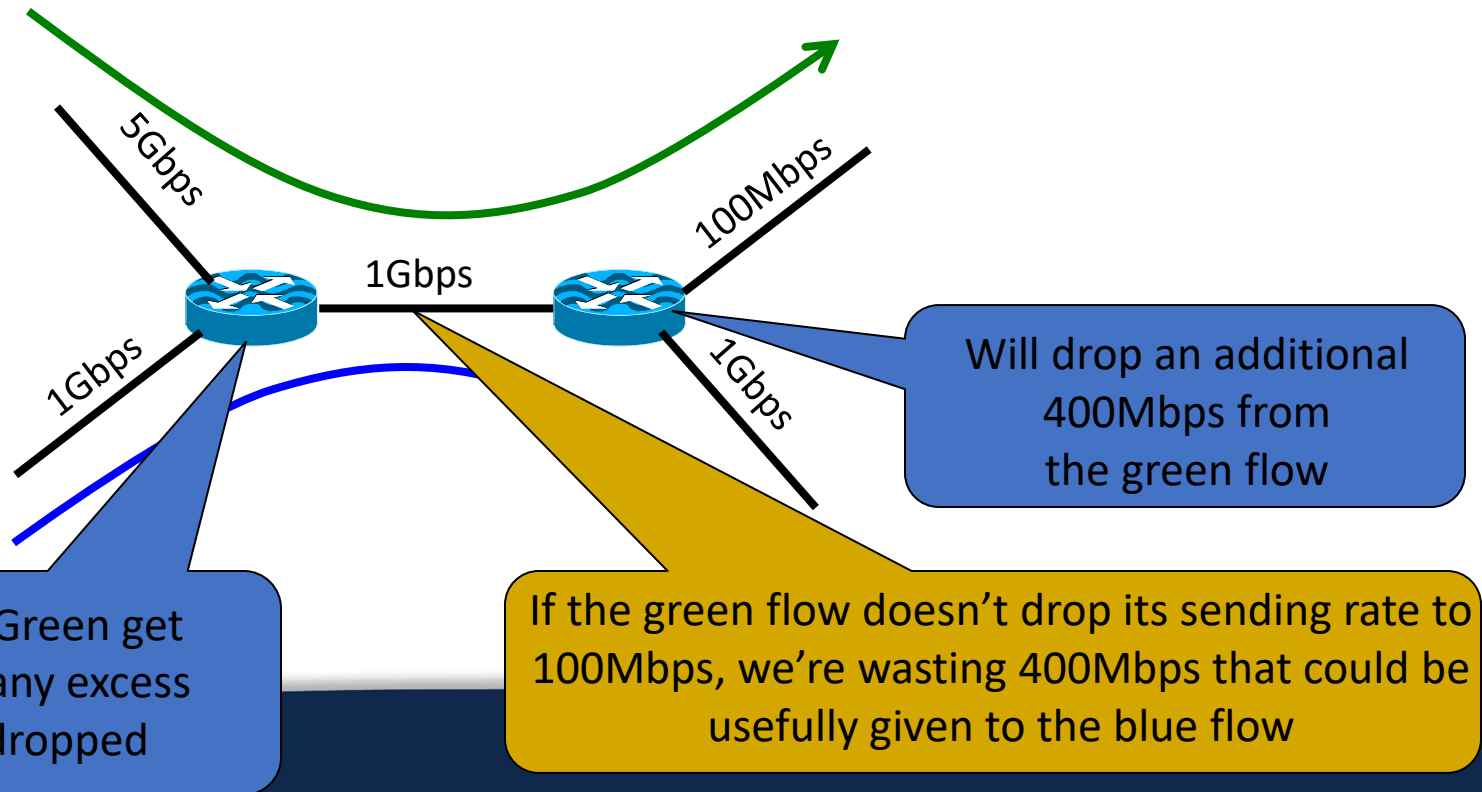
- Implementation of round-robin generalized to the case where not all packets are equal sized
- **Weighted fair queuing (WFQ)**: assign different flows different shares
- Today, some form of WFQ implemented in almost all routers
  - Not the case in the 1980-90s, when CC was being developed
  - Mostly used to isolate traffic at larger granularities (e.g., per-prefix)

# FQ vs. FIFO

- FQ advantages:
  - Isolation: cheating flows don't benefit
  - Bandwidth share does not depend on RTT
  - Flows can pick any rate adjustment scheme they want
- Disadvantages:
  - More complex than FIFO: per flow queue/state, additional per-packet book-keeping

# FQ in the big picture

- FQ does not eliminate congestion → it just manages the congestion



# FQ in the big picture

- FQ does not eliminate congestion → it just manages the congestion
  - Robust to cheating, variations in RTT, details of delay, reordering, retransmission, etc.
- But congestion (and packet drops) still occurs
- We still want end-hosts to discover/adapt to their fair share!
- What would the end-to-end argument say w.r.t. congestion control?

# Fairness is a controversial goal

- What if you have 8 flows, and I have 4?
  - Why should you get twice the bandwidth?
- What if your flow goes over 4 congested hops, and mine only goes over 1?
  - Why shouldn't you be penalized for using more scarce bandwidth?
- What is a flow anyway?
  - TCP connection
  - Source-Destination pair?
  - Source?

# Router-Assisted Congestion Control

- CC has three different tasks:
  - Isolation/fairness
  - Rate adjustment
  - Detecting congestion



# Why not let routers tell what rate end hosts should use?

- Packets carry “rate field”
- Routers insert “fair share”  $f$  in packet header
- End-hosts set sending rate (or window size) to  $f$ 
  - Hopefully (still need some policing of end hosts!)
- This is the basic idea behind the “Rate Control Protocol” (RCP) from Dukkupati et al. '07
  - Flows react faster

# Router-Assisted Congestion Control

- CC has three different tasks:
  - Isolation/fairness
  - Rate adjustment
  - Detecting congestion

# Explicit Congestion Notification (ECN)

- Single bit in packet header; set by congested routers
  - If data packet has bit set, then ACK has ECN bit set
- Many options for when routers set the bit
  - Tradeoff between (link) utilization and (packet) delay
- Congestion semantics can be exactly like that of drop
  - i.e., end-host reacts as though it saw a drop

# ECN

- Advantages:
  - Don't confuse corruption with congestion; recovery w/ rate adjustment
  - Can serve as an early indicator of congestion to avoid delays
  - Easy (easier) to incrementally deploy
    - Today: defined in RFC 3168 using ToS/DSCP bits in the IP header
    - Common in datacenters

# Routing

# Goal of routing

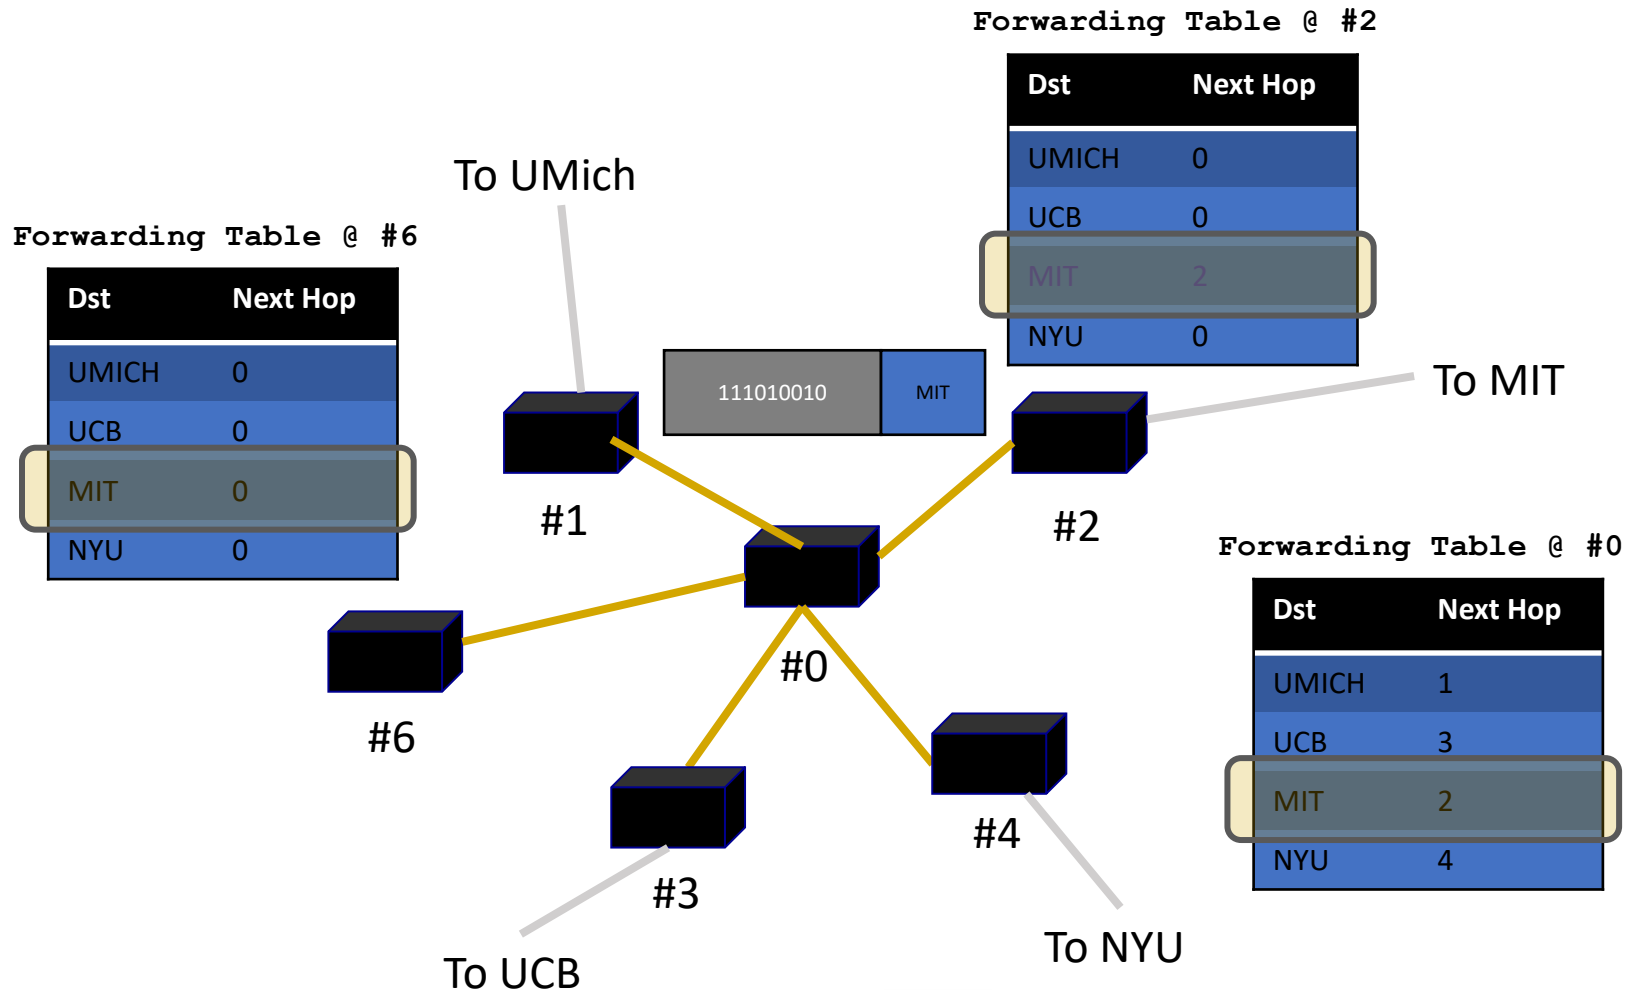
- Find a path to a given destination
- How do we know that the state contained in forwarding tables meets our goal?
  - This is what “[validity](#)” of routing state tells us
  - [\[This is non-standard terminology\]](#)

# Local vs. global view of state

- *Local* routing state is the forwarding table in a single router
  - By itself, the state in a single router cannot be evaluated
  - It must be evaluated in terms of the global context

# Example:

## Local vs. global view of state





# Local vs. global view of state

- *Local* routing state is the forwarding table in a single router
  - By itself, the state in a single router cannot be evaluated
  - It must be evaluated in terms of the global context
- *Global* state refers to the collection of forwarding tables in each of the routers
  - Global state determines which paths packets take
  - (Will discuss later where this routing state comes from)

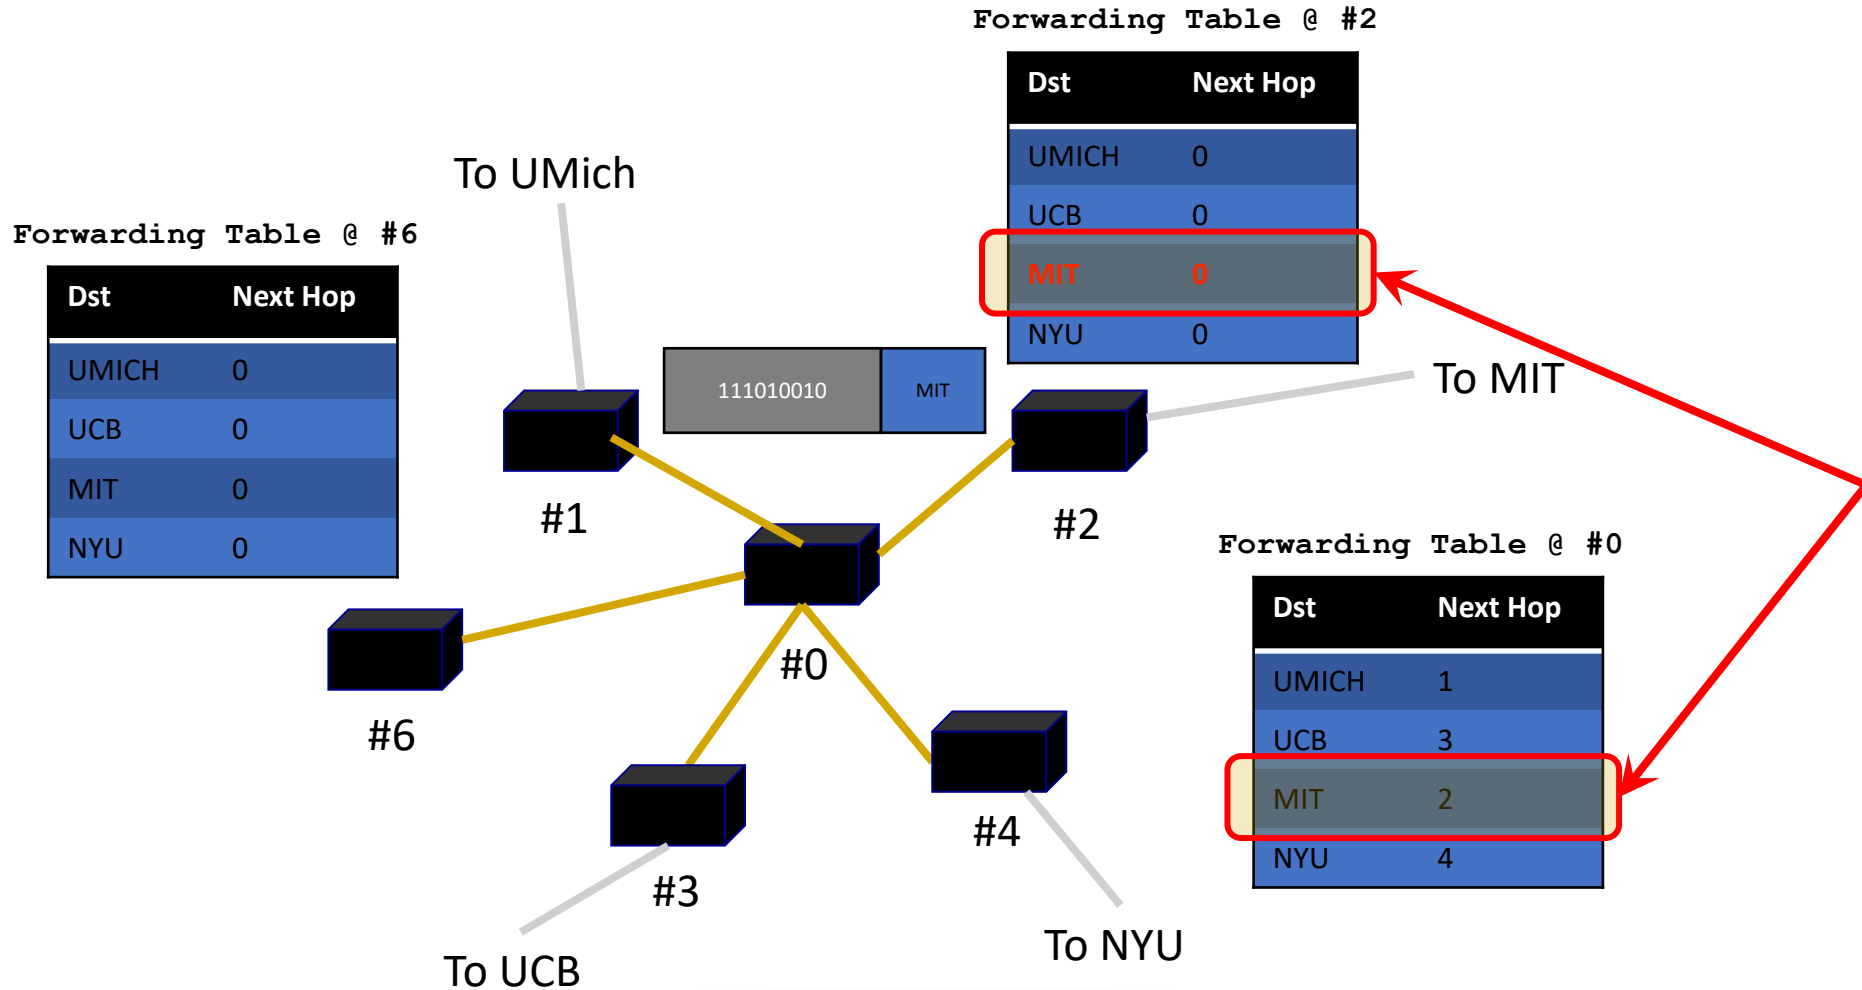
# “Valid” routing state

- Global state is “valid” if it produces forwarding decisions that always deliver packets to their destinations
- Goal of routing protocols: **compute valid state**
  - How can we tell if routing state is valid?
- Need a succinct correctness condition for routing

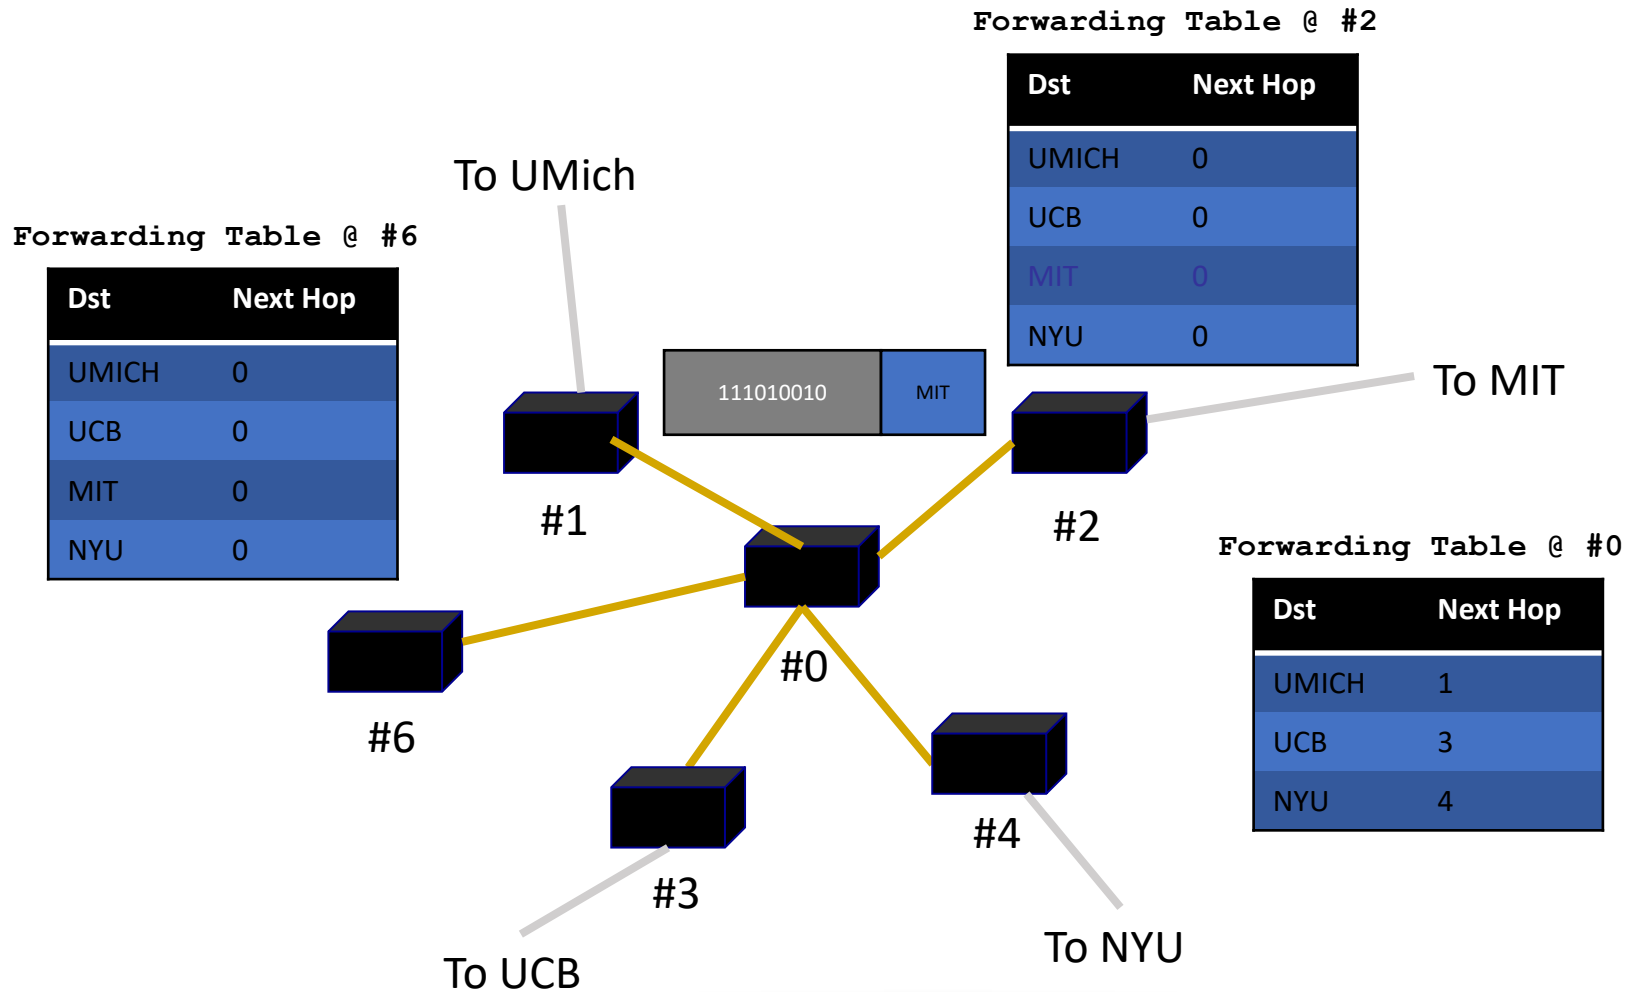
# Necessary and sufficient condition

- Global routing state is valid *if and only if*:
  - There are no dead ends (other than destination)
  - There are no loops
- A **dead end** is when there is no outgoing link (next-hop)
  - A packet arrives, but the forwarding decision does not yield any outgoing link
- A **loop** is when a packet cycles around the same set of nodes forever

# Loop!



# Dead end to MIT @ #0



# Necessary and sufficient condition

- Global routing state is valid *if and only if*:
  - There are no dead ends (other than destination)
  - There are no loops

# Necessary (“only if”)

- If you run into a dead end before hitting destination,
  - you’ll never reach the destination
- If you run into a loop,
  - you’ll never reach destination

# Sufficient (“if”)


- Assume there are no dead ends and no loops
- Packet must keep wandering, but without repeating
  - If ever enter same switch from same link, will loop
- Only a finite number of possible links for it to visit
  - It cannot keep wandering forever without looping
  - Must eventually hit destination




# Checking validity of routing state

- Focus only on a single destination
  - Ignore all other routing state
- Mark outgoing link (“next hop”) with arrow
  - There is only one at each node
- Eliminate all links with no arrows
- Look at what’s left


# Example 1




# Pick destination



Put arrows on outgoing links (to green dot)




# Remove unused links




Leaves spanning tree: Valid

## Example 2



Is this valid?

Not valid: Contains loop!



# Routing validity

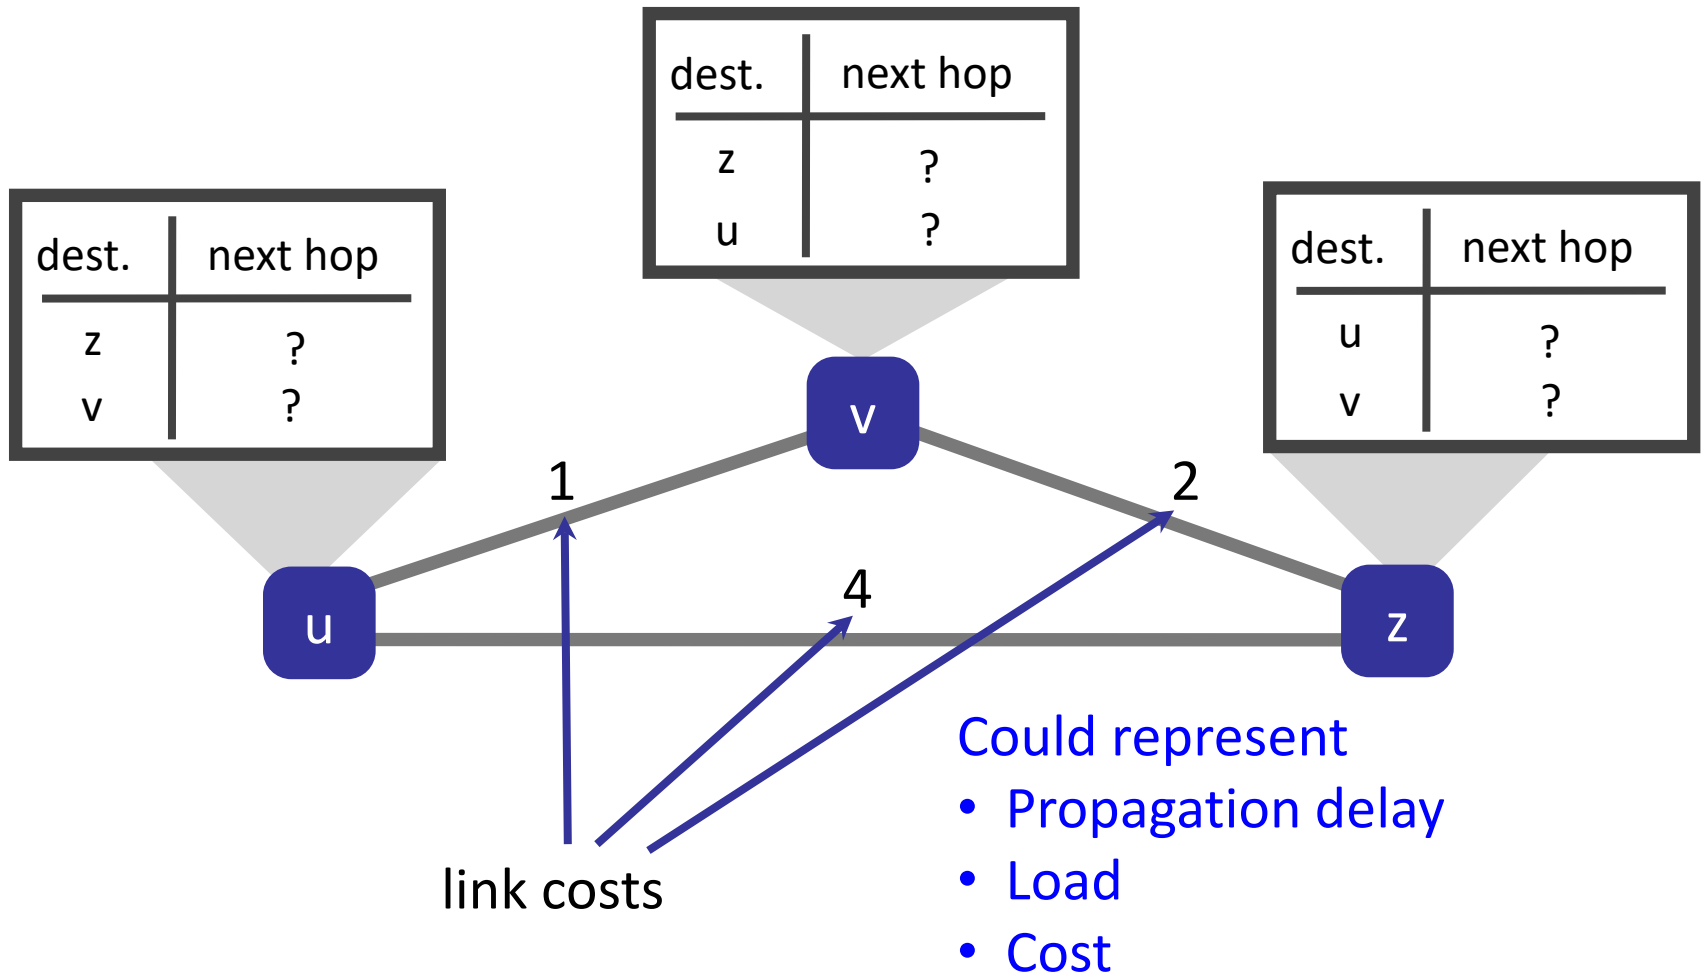
- Very easy to check validity of routing state for a particular destination
- Dead ends are nodes without outgoing arrow
- Loops are obvious too
  - Disconnected from rest of graph




# Goal of routing

- v1: Find a path to a given destination
- v2: Find a *least-cost path* to a given destination

# Example



# Example



least-cost path from u to z: u v z

least cost path from u to v: u v

# Least-cost path routing

- **Given:** router graph & link costs
- **Goal:** find least-cost path
  - From each source router to each destination router

# Least-cost routes

- Least-cost routes provide an easy way to avoid loops
  - No reasonable cost metric is minimized by traversing a loop
- Least-cost paths form a spanning tree for each destination rooted at that destination

# EECS 281:

## Dijkstra's algorithm

- Network topology, link costs known to all nodes
  - All nodes have same info
- Computes least-cost paths from one node ("src") to all other nodes
  - After  $k$  iterations, know least-cost path to  $k$  destinations

### • Notations

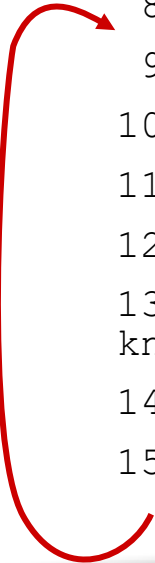
- $c(x,y)$ : link cost from  $x$  to  $y$ ;
  - $\infty$  if not direct neighbors
- $D(v)$ : current value of cost of path from src to dst  $v$
- $p(v)$ : predecessor node along path from source to  $v$
- $N'$ : set of nodes whose least-cost path definitively known

# Dijkstra's algorithm

```
1  Initialization:  
2     $N' = \{u\}; D(u) = 0$   
3    for all nodes  $v$   
4      if  $v$  adjacent to  $u$   
5        then  $D(v) = c(u, v)$   
6      else  $D(v) = \infty$ 
```

# Dijkstra's algorithm

```
1  Initialization:
2     $N' = \{u\}; D(u) = 0$ 
3    for all nodes  $v$ 
4        if  $v$  adjacent to  $u$ 
5            then  $D(v) = c(u, v)$ 
6            else  $D(v) = \infty$ 
7
8  Loop
9    find  $w$  not in  $N'$  such that  $D(w)$  is a minimum
10   add  $w$  to  $N'$ 
11   update  $D(v)$  for all  $v$  adjacent to  $w$  and not in  $N'$ :
12        $D(v) = \min( D(v), D(w) + c(w, v) )$ 
13       /* new cost to  $v$  is either old cost to  $v$  or
known
14         least path cost to  $w$  plus cost from  $w$  to  $v$  */
15  until all nodes are in  $N'$ 
```






# Dijkstra's algorithm: Example

Step	N'	D(v) p(v)	D(w) p(w)	D(x) p(x)	D(y) p(y)	D(z) p(z)
0	u	7,u	3,u	5,u	$\infty$	$\infty$
1	uw	6,w		5,u	11,w	$\infty$
2	uw x	6,w			11,w	14,x
3	uw x v			10,v		14,x
4	uw x v y				12,y	
5	uw x v y z					

## Notes:

- Construct shortest path tree by tracing predecessor nodes
- Ties can exist (can be broken arbitrarily)




# Dijkstra's algorithm: Example

*Resulting forwarding table  
in  $u$*

destination	link
v	(u, w)
w	(u, w)
x	(u, x)
y	(u, w)
z	(u, w)

*Resulting least-cost tree  
from  $u$*



# Let's Try to learn about the network

- Send a message to your neighbors with your id
- Build your routing table based on this information
- Send to your neighbors your routing table (without the interface)
- When you receive information about nodes you don't know of, update your table with your neighbor as next hop and send you table to your neighbors

Node ID – Interface – Hops		
5	1	1
8	2	1
4	3	1

Node ID – Interface – Hops		
7	2	1
15	3	1

Node ID – Interface – Hops		
5	1	1
8	2	1
4	3	1
7	1	3
15	1	4

# Let's Try to learn about the network

- Send a message to your neighbors
- Send a message to each neighbor, informing them about each one of your neighbors
- If you receive this message from a neighbor, make copies and forward one copy to each neighbor

# Summary

- Network layer control plane calculates valid routes and sets up forwarding table
  - Avoiding loops and dead ends
- Least-cost routes can be calculated using Dijkstra's algorithm
- [Next lecture](#): Routing protocols

# Bonus Quiz 11

- <https://forms.gle/8e6vJUgRfadYNvQK9>

