



EECS 489

Computer Networks

IntraAS Routing

Agenda

- Link-state routing
- Distance-vector routing

Recap: Least-cost path routing

- **Given:** router graph & link costs
- **Goal:** find least-cost path
 - From each source router to each destination router
- Easy way to avoid loops
 - No reasonable cost metric is minimized by traversing a loop

Recap:

Dijkstra's algorithm

- Network topology, link costs known to all nodes
 - All nodes have same info
- Each node (“src”) computes least-cost paths to all other nodes
 - After k iterations, know least-cost path to k destinations

From routing algorithm to protocol

- Dijkstra's is a local computation!
 - Computed by a node given complete network graph
- Possibilities:
 - Option#1: a separate machine runs the algorithm
 - Option#2: every router runs the algorithm
- The Internet currently uses Option#2

This is what we did last time

- Assigned node ids
- Round 1:
 - Say Hello to your neighbors:
 - someone sitting next to you
 - Someone sitting right behind you
 - Say “Hi” and tell them your node id
 - Write down the ids of your neighbors

Building the table

- Round 2:
 - Made a copy of your neighbors and sent it to each one of your neighbors
 - For example, I am node 1 and my neighbors are 2, 4, and 5, I sent the following to my neighbors

node – hops – nexthop		
4	1	4
5	1	5
2	1	2

Updates

- If you receive cost to a node that is less than your current cost, add that to your routing table:

Your routing table is:

Node	hops	nexthop
4	1	4
5	1	5
2	1	2
8	5	2

You receive from node 5 the following information:

Node	hops	nexthop
3	1	3
1	1	1
8	2	7

Update your table to the following and send it to your neighbors

Node	hops	nexthop
4	1	4
5	1	5
2	1	2
8	3	5
3	2	5

Any observations?

Let's try a different approach

- Assign node ids
- Round 1:
 - Say Hello to your neighbors:
 - someone sitting next to you
 - Someone sitting right behind you (could be a few rows behind)
 - Say “Hi” and tell them your node id
 - Write down the ids of your neighbors

Building the table

- Round 2:
 - Make a copy of your neighbors and send it to each one of your neighbors
 - For example, I am node 1 and my neighbors are 2, 4, and 5, I will send the following to my neighbors

node – hops – nexthop		
4	1	4
5	1	5
2	1	2

- when you receive a table from your neighbor, make a copy of it and send to all your neighbors – except the one who sent it to you

Network Graph

- Based on the tables that you have, can you draw a network graph?
- If yes, then you can use Dijkstra's algorithm to find best path to each node

Any observations?

Link-state routing

- Every router knows its local “link state”
 - Router u: “(u,v) with cost=2; (u,x) with cost=1”
- Each router floods its **local link state to all other routers** in the network
 - Does so periodically or when its link state changes
- Every router learns the entire network graph
 - Each runs Dijkstra’s Shortest-Path First (SPF) algorithm locally to compute forwarding table

Flooding link state

- Flooding
 - A node sends its link-state info out on all of its links
 - The next node forwards the info on all of its links except the one the information arrived at
- When to initiate flooding?
 - Topology change (e.g., link/node failure/recovery)
 - Configuration change (e.g., link cost change)
 - Periodically
 - To refresh link-state information (soft states)
 - Typically (say) every 30 minutes


Convergence

- Why flood?
 - To get all the nodes in the network to **converge** to the new topology
- Upon convergence, all nodes will have **consistent routing information** and can **compute consistent forwarding**:
 - All nodes have the same link-state database
 - All nodes forward packets on shortest paths
 - The next router on the path forwards to the expected next hop


Convergence delay

- Time to achieve convergence
- Sources of convergence delay
 - Time to detect failure
 - Time to flood link-state information
 - Time to re-compute forwarding tables
- What happens if it takes too long to converge?

Loop from convergence delay



u and w think that the path to y goes through v



v thinks that the path to y goes through w

Performance during convergence period

- Looping packets
- Lost packets due to black holes
- Out-of-order packets reaching the destination

Link-state routing

- Scalability?
 - $O(NE)$ messages
 - $O(N^2)$ computation time
 - $O(\text{Network diameter})$ convergence delay
 - $O(N)$ entries in forwarding table

Link-state routing protocols

- **OSPF**: Open Shortest Path First
- **IS-IS**: Intermediate System to Intermediate System
 - Similar to OSPF

OSPF:

Open Shortest-Path First

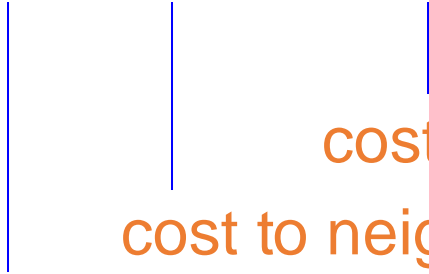
- **Open**: publicly available
- Uses link-state algorithm
 - Link-state packet dissemination
 - Topology map at each node
 - Route computation using Dijkstra's algorithm
- Router floods OSPF link-state advertisements to all other routers in entire AS
 - Carried in OSPF messages **directly over IP** (rather than TCP or UDP)
 - Requires reliable transmission

Distance-vector protocol


- Link-state routing protocol
 - Each node **broadcasts** its **local** information
- Distance-vector routing protocol
 - The opposite (sort of)
 - Each node **tells its neighbors** about its **global** view

Bellman-Ford equation

- Let
 - $d_x(y) := \text{cost of least-cost path from } x \text{ to } y$
- Then
 - $d_x(y) = \min_v \{c(x, v) + d_v(y)\}$


cost from neighbor v to destination y
cost to neighbor v
min taken over all neighbors v of x

Bellman-Ford example



$$d_x(z) = 9, d_w(z) = 9, d_v(z) = 6$$

B-F equation says:

$$\begin{aligned} d_u(z) &= \min \{c(u,x) + d_x(z), \\ &\quad c(u,w) + d_w(z), \\ &\quad c(u,v) + d_v(z) \} \\ &= \min \{5 + 9, \\ &\quad 3 + 9, \\ &\quad 7 + 6\} = 12 \end{aligned}$$

Neighbor achieving the minimum (w) is next hop in shortest path, used in forwarding table

Distance vector algorithm

- $D_x(y)$ is the estimate of least cost from x to y
 - x maintains its own distance vector $\mathbf{D}_x = [D_x(y): y \in N]$
- Node x :
 - Knows cost to each neighbor v : $c(x,v)$
 - Maintains its neighbors' distance vectors
 - For each neighbor v , x has $\mathbf{D}_v = [D_v(y): y \in N]$

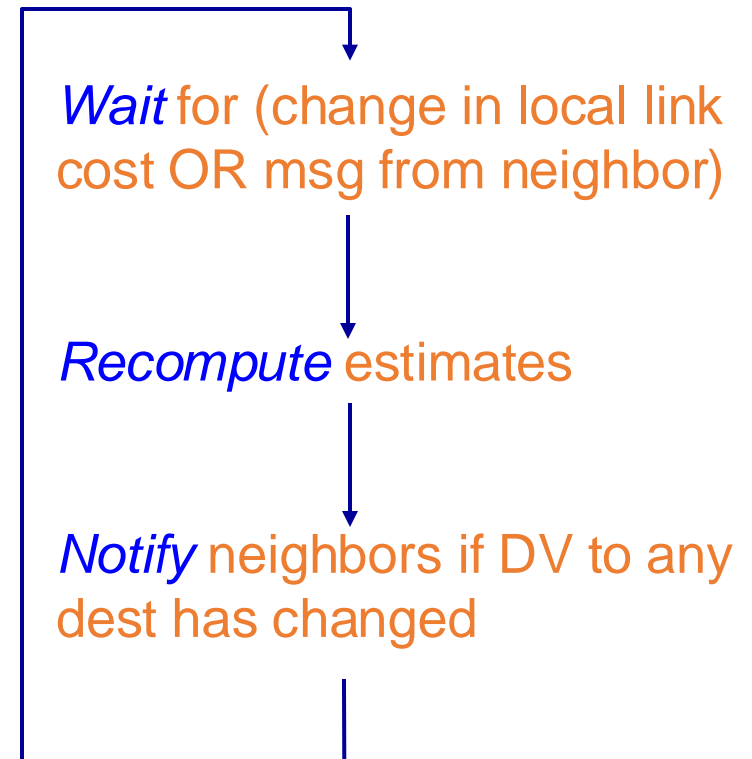
Distance vector algorithm

- From time-to-time, each node sends its own distance vector estimate to neighbors
- When x receives new DV estimate from neighbor, it updates its own DV using B-F equation
 - $D_x(y) \leftarrow \min_v \{c(x,v) + D_v(y)\}$ for each node $y \in N$
- Eventually, the estimate $D_x(y)$ may converge to the actual least cost $d_x(y)$

Distance vector algorithm


- **Iterative, asynchronous**
 - Local iterations caused by
 - Local link cost change
 - DV update message from neighbor
- **Distributed**
 - Each node notifies neighbors only when its DV changes
 - Neighbors then notify their neighbors if necessary

@each node:



Example

	x	y	z
y	2	0	1
z	3	1	0




	x	y	z
y	2	0	1
z	3	1	0

Example

	x	y	z
y	4	0	1
z	3	1	0

routing loop!




	x	y	z
y	2	0	1
z	3	1	0

Example

	x	y	z
y	4	0	1
z	3	1	0

routing loop!




	x	y	z
y	4	0	1
z	3	1	0

Example

	x	y	z
y	4	0	1
z	3	1	0

routing loop!




	x	y	z
y	4	0	1
z	5	1	0

Example

	x	y	z
y	4	0	1
z	5	1	0

routing loop!




	x	y	z
y	4	0	1
z	5	1	0

Example

	x	y	z
y	6	0	1
z	5	1	0

routing loop!




	x	y	z
y	4	0	1
z	5	1	0

Example

	x	y	z
y	6	0	1
z	5	1	0

routing loop!




	x	y	z
y	6	0	1
z	5	1	0

Example

	x	y	z
y	6	0	1
z	5	1	0

routing loop!




	x	y	z
y	6	0	1
z	7	1	0

Example

	x	y	z
y	6	0	1
z	7	1	0

routing loop!




	x	y	z
y	6	0	1
z	7	1	0

Example

	x	y	z
y	8	0	1
z	7	1	0

routing loop!




	x	y	z
y	6	0	1
z	7	1	0

Example

	x	y	z
y	8	0	1
z	7	1	0

routing loop!




	x	y	z
y	8	0	1
z	7	1	0

Example

	x	y	z
y	8	0	1
z	7	1	0

routing loop!



Count-to-infinity scenario

	x	y	z
y	8	0	1
z	7	1	0


Problems with Bellman-Ford

- Routing loops
 - z routes through y, y routes through x
 - y loses connectivity to x
 - y decides to route through z
- Can take a very long time to resolve
 - Count-to-infinity scenario

Poisoned reverse


- One **heuristic** to avoid count-to-infinity
 - If z routes to x through y,
 - z advertises to y that its cost to x is infinite
 - y never decides to route to x through z

	x	y	z
y	2	0	1
z	∞	1	0




	x	y	z
y	2	0	1
z	3	1	0

	x	y	z
y	∞	0	1
z	∞	1	0




	x	y	z
y	2	0	1
z	3	1	0

	x	y	z
y	∞	0	1
z	∞	1	0




	x	y	z
y	∞	0	1
z	3	1	0

	x	y	z
y	∞	0	1
z	∞	1	0




	x	y	z
y	∞	0	1
z	7	1	0

	x	y	z
y	∞	0	1
z	7	1	0



	x	y	z
y	∞	0	1
z	7	1	0

	x	y	z
y	8	0	1
z	7	1	0



	x	y	z
y	∞	0	1
z	7	1	0

Poisoned reverse

- One **heuristic** to avoid count-to-infinity
 - If z routes to x through y,
 - z advertises to y that its cost to x is infinite
 - y never decides to route to x through z
- **Not guaranteed**
- **Loop-free routing** examples include
 - Path vector
 - Source tracing

Distance-vector routing

- Scalability?
 - Requires fewer messages than Link-State
 - $O(N)$ update time on arrival of a new DV from neighbor
 - $O(\text{network diameter})$ convergence time
 - $O(N)$ entries in forwarding table
- RIP is a protocol that implements DV (IETF RFC 2080)

Comparison of LS and DV routing

Messaging complexity

- LS: with N nodes, E links, $O(NE)$ messages sent
- DV: exchange between neighbors only

Speed of convergence

- LS: relatively fast
- DV: convergence time varies
 - Count-to-infinity problem

Robustness: what happens if router malfunctions?

- LS:
 - Node can advertise incorrect [link](#) cost
 - Each node computes its *own* table
- DV:
 - Node can advertise incorrect [path](#) cost
 - Each node's table used by others (error propagates)

Similarities between LS and DV routing

- Both are shortest-path based routing
 - Minimizing cost metric (link weights) a common optimization goal
 - Routers share a common view as to what makes a path “good” and how to measure the “goodness” of a path
- Due to shared goal, commonly used inside an organization
 - RIP and OSPF are mostly used for **intra**-domain routing

Summary

- Intra-AS routing
 - Link-state routing
 - Distance-vector routing


Bonus Quiz 12

- <https://forms.gle/QgMjvaz5fLNyTAWR6>



Backup

- Poison reverse doesn't always work.



- Link between C and D breaks
 - C will try to get to D through B that gets to D through A